

```
In [1]: pip install -U scikit-learn imbalanced-learn
```

```
Requirement already satisfied: scikit-learn in c:\users\kariu\anaconda3\lib\site-packages (1.3.2)
Requirement already satisfied: imbalanced-learn in c:\users\kariu\appdata\roaming\python\python311\site-packages (0.11.0)
Requirement already satisfied: numpy<2.0,>=1.17.3 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (1.24.3)
Requirement already satisfied: scipy>=1.5.0 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (1.11.1)
Requirement already satisfied: joblib>=1.1.1 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (2.2.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [2]: pip install --user -U scikit-learn imbalanced-learn
```

```
Requirement already satisfied: scikit-learn in c:\users\kariu\anaconda3\lib\site-packages (1.3.2)
Requirement already satisfied: imbalanced-learn in c:\users\kariu\appdata\roaming\python\python311\site-packages (0.11.0)
Requirement already satisfied: numpy<2.0,>=1.17.3 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (1.24.3)
Requirement already satisfied: scipy>=1.5.0 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (1.11.1)
Requirement already satisfied: joblib>=1.1.1 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\kariu\anaconda3\lib\site-packages (from scikit-learn) (2.2.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [88]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
```

```
from datetime import datetime
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve
from sklearn.ensemble import RandomForestClassifier
from imblearn.over_sampling import SMOTE
from imblearn.combine import SMOTEENN
from scipy.stats import mannwhitneyu
from sklearn.preprocessing import StandardScaler
from sklearn.utils.class_weight import compute_class_weight
from sklearn.cluster import KMeans
from tqdm.notebook import tqdm
from sklearn.pipeline import make_pipeline
```

```
In [4]: client_data = pd.read_csv("client_data.csv")
price_data = pd.read_csv("price_data.csv")
```

Data types

```
In [5]: price_data.dtypes
```

```
Out[5]: id                object
price_date              object
price_off_peak_var      float64
price_peak_var          float64
```

```
price_mid_peak_var      float64
price_off_peak_fix      float64
price_peak_fix          float64
price_mid_peak_fix      float64
dtype: object
```

```
In [6]: client_data.dtypes
```

```
Out[6]: id                object
channel_sales            object
cons_12m                 int64
cons_gas_12m             int64
cons_last_month          int64
date_activ              object
date_end                object
date_modif_prod          object
date_renewal            object
forecast_cons_12m        float64
forecast_cons_year       int64
forecast_discount_energy float64
forecast_meter_rent_12m  float64
forecast_price_energy_off_peak float64
forecast_price_energy_peak float64
forecast_price_pow_off_peak float64
has_gas                 object
imp_cons               float64
margin_gross_pow_ele    float64
margin_net_pow_ele      float64
nb_prod_act             int64
net_margin              float64
num_years_antig         int64
origin_up              object
pow_max                float64
churn                  int64
dtype: object
```

Descriptive statistics

Information

```
In [7]: client_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    14606 non-null  object
1   channel_sales                        14606 non-null  object
2   cons_12m                             14606 non-null  int64
3   cons_gas_12m                         14606 non-null  int64
4   cons_last_month                      14606 non-null  int64
5   date_activ                           14606 non-null  object
6   date_end                             14606 non-null  object
7   date_modif_prod                      14606 non-null  object
8   date_renewal                         14606 non-null  object
9   forecast_cons_12m                    14606 non-null  float64
10  forecast_cons_year                    14606 non-null  int64
11  forecast_discount_energy              14606 non-null  float64
12  forecast_meter_rent_12m               14606 non-null  float64
13  forecast_price_energy_off_peak        14606 non-null  float64
14  forecast_price_energy_peak            14606 non-null  float64
15  forecast_price_pow_off_peak           14606 non-null  float64
16  has_gas                               14606 non-null  object
17  imp_cons                             14606 non-null  float64
```

```
18 margin_gross_pow_ele      14606 non-null float64
19 margin_net_pow_ele        14606 non-null float64
20 nb_prod_act                14606 non-null int64
21 net_margin                 14606 non-null float64
22 num_years_antig            14606 non-null int64
23 origin_up                  14606 non-null object
24 pow_max                    14606 non-null float64
25 churn                      14606 non-null int64
dtypes: float64(11), int64(7), object(8)
memory usage: 2.9+ MB
```

```
In [8]: price_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     193002 non-null object
1   price_date             193002 non-null object
2   price_off_peak_var     193002 non-null float64
3   price_peak_var         193002 non-null float64
4   price_mid_peak_var     193002 non-null float64
5   price_off_peak_fix     193002 non-null float64
6   price_peak_fix         193002 non-null float64
7   price_mid_peak_fix     193002 non-null float64
dtypes: float64(6), object(2)
memory usage: 11.8+ MB
```

Summary

```
In [9]: client_data.describe()
```

Out[9]:

	cons_12m	cons_gas_12m	cons_last_month	forecast_cons_12m	forecast_cons_year	forecast_discount_ene
count	1.460600e+04	1.460600e+04	14606.000000	14606.000000	14606.000000	14606.000
mean	1.592203e+05	2.809238e+04	16090.269752	1868.614880	1399.762906	0.966
std	5.734653e+05	1.629731e+05	64364.196422	2387.571531	3247.786255	5.108
min	0.000000e+00	0.000000e+00	0.000000	0.000000	0.000000	0.000
25%	5.674750e+03	0.000000e+00	0.000000	494.995000	0.000000	0.000
50%	1.411550e+04	0.000000e+00	792.500000	1112.875000	314.000000	0.000
75%	4.076375e+04	0.000000e+00	3383.000000	2401.790000	1745.750000	0.000
max	6.207104e+06	4.154590e+06	771203.000000	82902.830000	175375.000000	30.000

```
In [10]: price_data.describe()
```


Out[10]:

	price_off_peak_var	price_peak_var	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price_mid_peak_fix
count	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000
mean	0.141027	0.054630	0.030496	43.334477	10.622875	6.409984
std	0.025032	0.049924	0.036298	5.410297	12.841895	7.773592
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.125976	0.000000	0.000000	40.728885	0.000000	0.000000
50%	0.146033	0.085483	0.000000	44.266930	0.000000	0.000000
75%	0.151635	0.101673	0.072558	44.444710	24.339581	16.226385

max 0.280700 0.229788 0.114102 59.444710 36.490692 17.458221


```
In [12]: client_data_numeric_columns = client_data.select_dtypes(include=['float64', 'int64'])

plt.figure(figsize=(12, 8))
sns.heatmap(client_data_numeric_columns.corr(), annot=True, cmap='coolwarm', fmt='.2f',
plt.title('Correlation Heatmap - Client Data')
plt.show()
```



```
In [13]: price_data_numeric_columns = price_data.select_dtypes(include=['float64', 'int64'])

plt.figure(figsize=(12, 8))
sns.heatmap(price_data_numeric_columns.corr(), annot=True, cmap='coolwarm', fmt='.2f', 1
plt.title('Correlation Heatmap - Client Data')
plt.show()
```




Distribution of Client Data

```
In [14]: sns.set(style="whitegrid")

numeric_columns = client_data.select_dtypes(include=['float64', 'int64'])

# Create a pairplot
sns.pairplot(numeric_columns)
plt.show()
```

C:\Users\kariu\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)



```
In [15]: # Set the style of seaborn
sns.set(style='whitegrid')

# Plotting distributions of a few numerical columns
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(14, 10))

# Distribution of electricity consumption over 12 months
sns.histplot(client_data['cons_12m'], bins=50, ax=axes[0, 0], kde=True)
axes[0, 0].set_title('Distribution of Electricity Consumption Over 12 Months')
axes[0, 0].set_xscale('log')

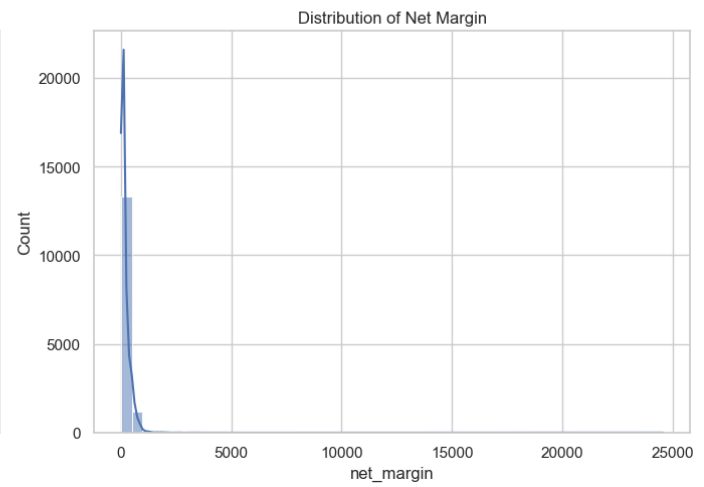

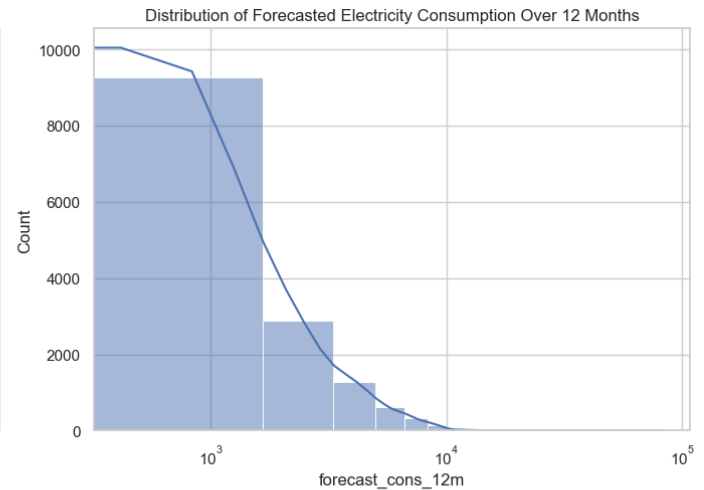
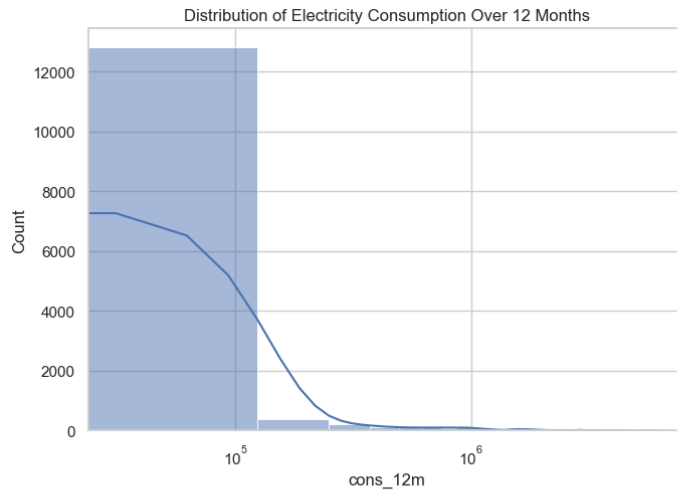
# Distribution of forecasted electricity consumption over 12 months
sns.histplot(client_data['forecast_cons_12m'], bins=50, ax=axes[0, 1], kde=True)
axes[0, 1].set_title('Distribution of Forecasted Electricity Consumption Over 12 Months')
axes[0, 1].set_xscale('log')

# Distribution of gross margin on power
sns.histplot(client_data['margin_gross_pow_ele'], bins=50, ax=axes[1, 0], kde=True)
axes[1, 0].set_title('Distribution of Gross Margin on Power')

# Distribution of net margin
```


```
sns.histplot(client_data['net_margin'], bins=50, ax=axes[1, 1], kde=True)
axes[1, 1].set_title('Distribution of Net Margin')

plt.tight_layout()
plt.show()
```



```
In [16]: sns.set(style='whitegrid')

# Plotting the churn distribution
plt.figure(figsize=(6, 4))
sns.countplot(x='churn', data=client_data)
plt.title('Churn Distribution')
plt.xlabel('Churn')
plt.ylabel('Number of Clients')
plt.xticks([0, 1], ['No', 'Yes'])
plt.show()
```

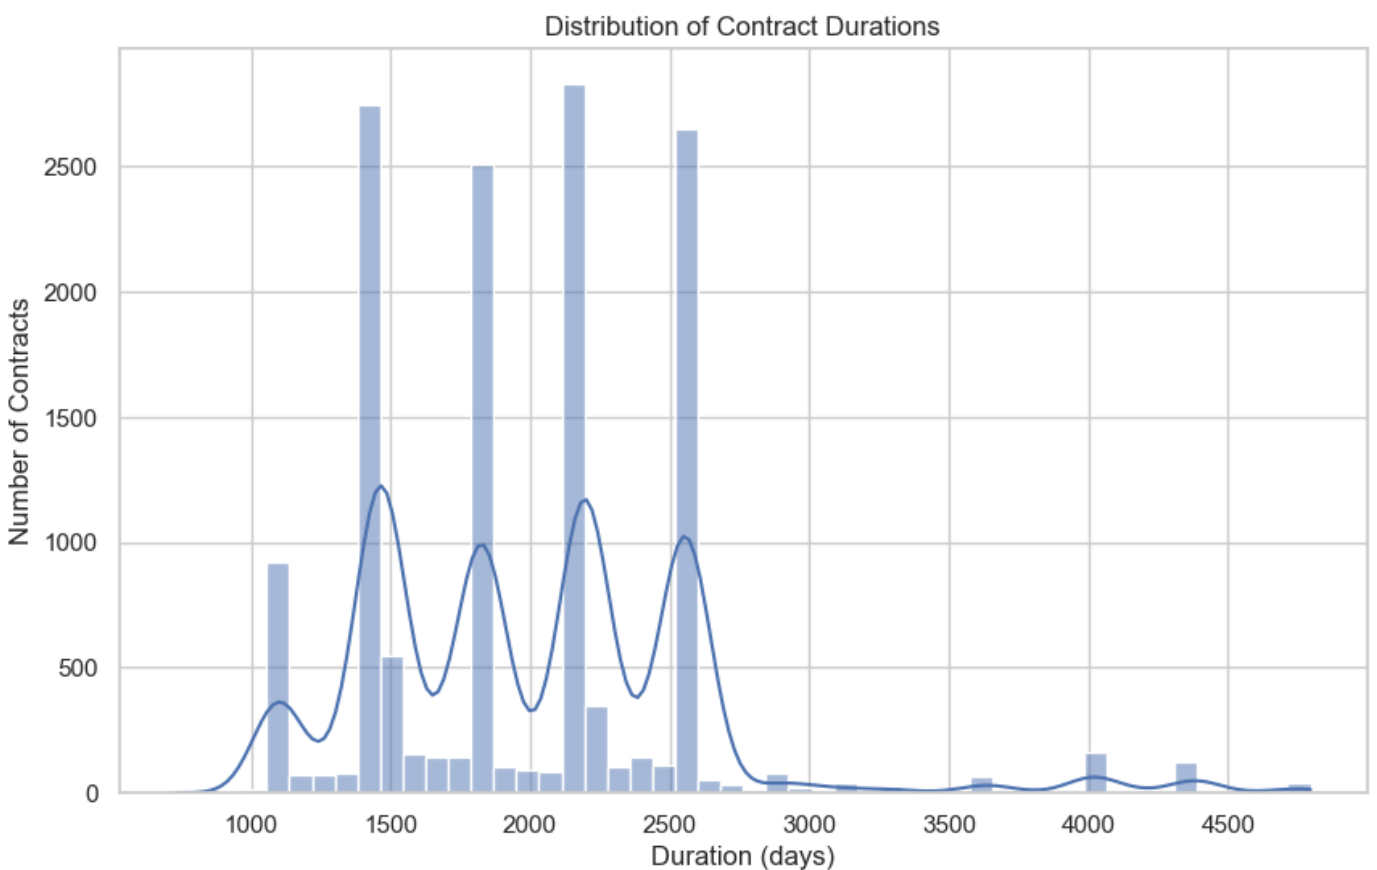


```
In [17]: # Convert date columns to datetime
client_data['date_activ'] = pd.to_datetime(client_data['date_activ'], format='%Y-%m-%d')
client_data['date_end'] = pd.to_datetime(client_data['date_end'], format='%Y-%m-%d', err
client_data['date_modif_prod'] = pd.to_datetime(client_data['date_modif_prod'], format='
client_data['date_renewal'] = pd.to_datetime(client_data['date_renewal'], format='%Y-%m-

# Calculate the duration of the contract in days
client_data['contract_duration'] = (client_data['date_end'] - client_data['date_activ'])

# Replace negative and NaN values with the median duration
median_duration = client_data['contract_duration'].median()
client_data['contract_duration'] = client_data['contract_duration'].apply(lambda x: medi

# Show the distribution of contract durations
plt.figure(figsize=(10, 6))
sns.histplot(client_data['contract_duration'], bins=50, kde=True)
plt.title('Distribution of Contract Durations')
plt.xlabel('Duration (days)')
plt.ylabel('Number of Contracts')
plt.show()
```

```
In [18]: # Calculate the average consumption and margins
client_data['average_consumption'] = client_data[['cons_12m', 'cons_gas_12m', 'cons_last
client_data['average_margin'] = client_data[['margin_gross_pow_ele', 'margin_net_pow_ele

# Calculate the average forecast consumption and price
client_data['average_forecast_cons'] = client_data[['forecast_cons_12m', 'forecast_cons
client_data['average_forecast_price'] = client_data[['forecast_price_energy_off_peak', '

# Summary statistics for the new average columns
summary_statistics = client_data[['average_consumption', 'average_margin', 'average_fore


# Display the summary statistics
summary_statistics
```

```
Out[18]:
```

	average_consumption	average_margin	average_forecast_cons	average_forecast_price
count	1.460600e+04	14606.000000	14606.000000	14606.000000
mean	6.780098e+04	24.563819	1634.188893	14.439277
std	2.433520e+05	20.230291	2563.882762	1.496629
min	0.000000e+00	0.000000	0.000000	0.000000
25%	2.225750e+03	14.280000	334.300000	13.608061
50%	5.814667e+03	21.640000	839.617500	14.818318
75%	1.797083e+04	29.880000	2026.855000	14.852151
max	2.408666e+06	374.640000	97280.390000	19.846780

```
In [19]: # Correlation matrix of the newly calculated averages and churn
# Selecting relevant columns for correlation
columns_of_interest = ['average_consumption', 'average_margin', 'average_forecast_cons',
correlation_matrix = client_data[columns_of_interest].corr()
```

```
# Plotting the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm')
plt.title('Correlation Matrix of Averages and Churn')
plt.show()
```



```
In [20]: # Preparing the data for logistic regression
X = client_data[['average_consumption', 'average_margin', 'average_forecast_cons', 'average_forecast_price']]
y = client_data['churn']

# Splitting the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initializing the logistic regression model
log_reg = LogisticRegression()

# Fitting the model
log_reg.fit(X_train, y_train)

# Predicting on the test set
y_pred = log_reg.predict(X_test)

# Generating the classification report and confusion matrix
report = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
# Output the classification report and confusion matrix
print('Classification Report:\n', report)
print('Confusion Matrix:\n', conf_matrix)
```

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.99	0.94	2617
1	0.04	0.00	0.01	305
accuracy			0.89	2922
macro avg	0.47	0.50	0.47	2922
weighted avg	0.81	0.89	0.84	2922

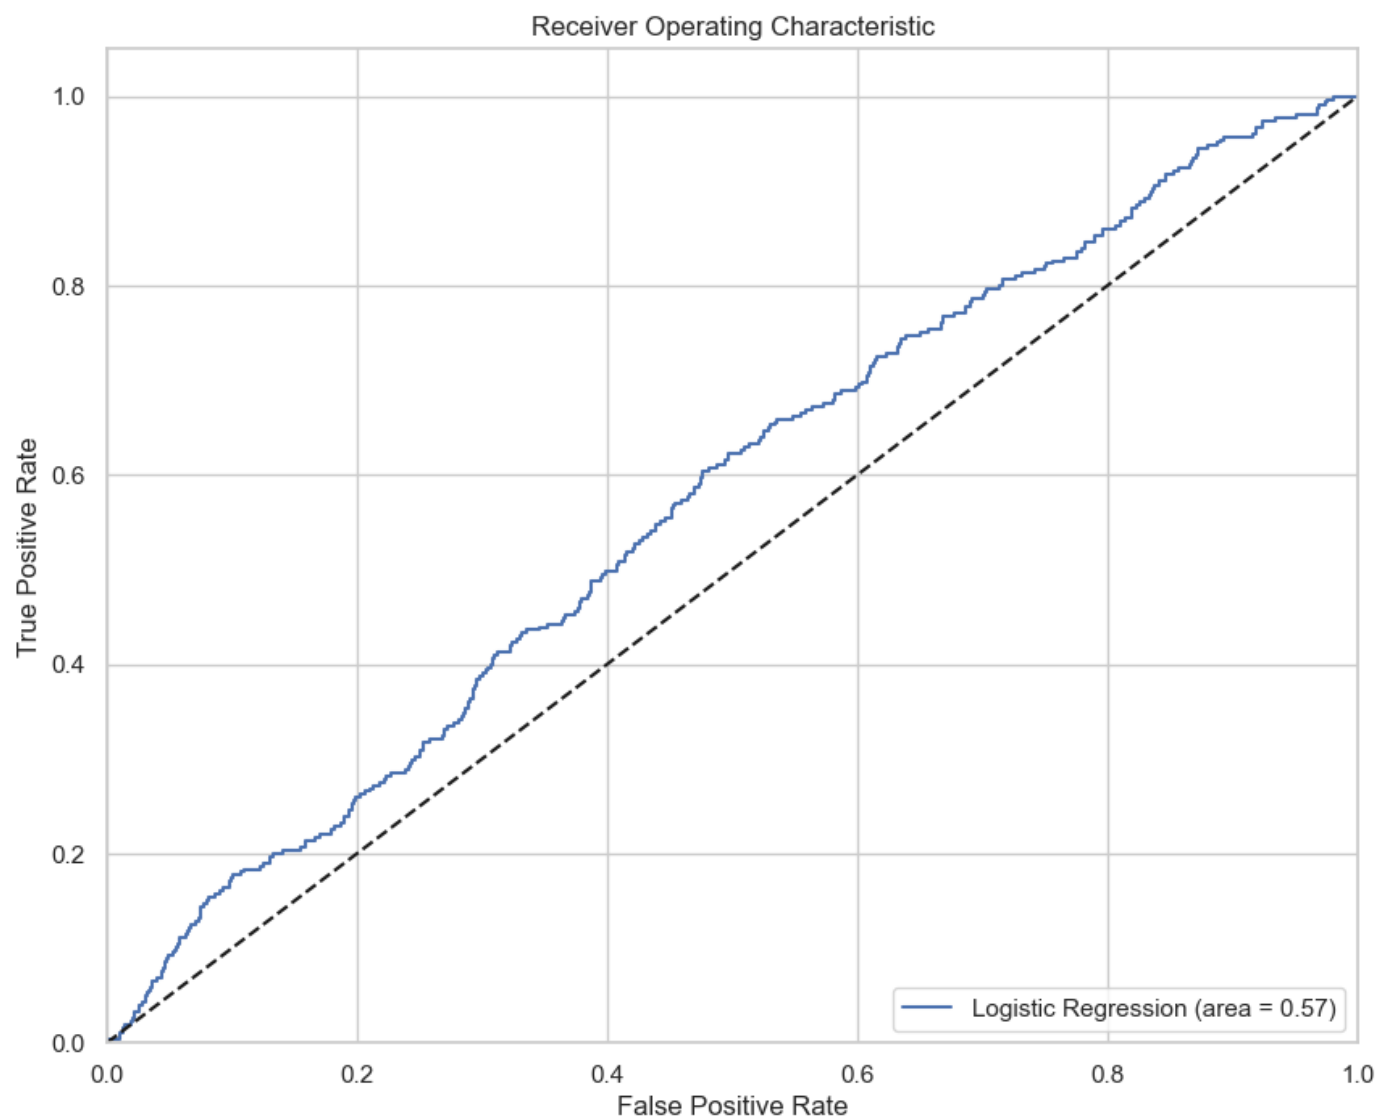
Confusion Matrix:

```
[[2595  22]
 [ 304   1]]
```

```
In [21]: # Calculate the ROC AUC score
roc_auc = roc_auc_score(y_test, log_reg.predict_proba(X_test)[:, 1])

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, log_reg.predict_proba(X_test)[:, 1])

# Plotting the ROC curve
plt.figure(figsize=(10, 8))
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
In [22]: # Define a parameter grid to search for the best parameters for random forest
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : [4,5,6,7,8],
    'criterion' :['gini', 'entropy']
}

# Create a base model
rf = RandomForestClassifier(random_state=42)

# Instantiate the grid search model
grid_search = GridSearchCV(estimator = rf, param_grid = param_grid,
                           cv = 3, n_jobs = -1, verbose = 2, scoring = 'roc_auc')

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Best parameters
best_params = grid_search.best_params_

# Output the best parameters
print('Best parameters found:\n', best_params)
```

Fitting 3 folds for each of 90 candidates, totalling 270 fits

C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection_validation.py:425: FitFailedWarning:
90 fits failed out of a total of 270.
Thescore on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting `error_score='raise'`.

Below are more details about the failures:

43 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py", line 729, in _fit_and_score
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\base.py", line 1145, in wrapper
```

```
    estimator._validate_params()
```

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\base.py", line 638, in _validate_params
```

```
    validate_parameter_constraints(
```

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\utils\_param_validation.py", line 96, in validate_parameter_constraints
```

```
    raise InvalidParameterError(
```

```
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.
```

47 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py", line 729, in _fit_and_score
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\base.py", line 1145, in wrapper
```

```
    estimator._validate_params()
```

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\base.py", line 638, in _validate_params
```

```
    validate_parameter_constraints(
```

```
File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\utils\_param_validation.py", line 96, in validate_parameter_constraints
```

```
    raise InvalidParameterError(
```

```
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
```

```
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:979: UserWarning: One or more of the test scores are non-finite: [      nan      nan      nan
0.64439381 0.64362271 0.64146143
```

```
0.64439381 0.64362271 0.64146143      nan      nan      nan
```

```
0.6477769  0.64773825 0.64840831 0.6477769  0.64773825 0.64840831
```

```
      nan      nan      nan 0.64869002 0.65039817 0.6490806
```

```
0.64869002 0.65039817 0.6490806      nan      nan      nan
```

```
0.65144331 0.65251391 0.65299086 0.65144331 0.65251391 0.65299086
```

```
      nan      nan      nan 0.64932779 0.65320376 0.65389176
```

```
0.64932779 0.65320376 0.65389176      nan      nan      nan
```

```
0.64282174 0.64277067 0.64148863 0.64282174 0.64277067 0.64148863
```

```
      nan      nan      nan 0.64412889 0.64613738 0.64555232
```

```
0.64412889 0.64613738 0.64555232      nan      nan      nan
```

```
0.64937615 0.65093337 0.65100836 0.64937615 0.65093337 0.65100836
```

```
      nan      nan      nan 0.65452613 0.6531797  0.65333803
```

```
0.65452613 0.6531797  0.65333803      nan      nan      nan
```

```
0.65509807 0.65660971 0.65769386 0.65509807 0.65660971 0.65769386]
```

```
warnings.warn(
```

Best parameters found:

```
{'criterion': 'entropy', 'max_depth': 8, 'max_features': 'sqrt', 'n_estimators': 200}
```

```

# 'auto' is equivalent to 'sqrt', so we will use 'sqrt' instead
rf_best_corrected = RandomForestClassifier(criterion='entropy', max_depth=8, max_feature

# Fitting the model
rf_best_corrected.fit(X_train, y_train)

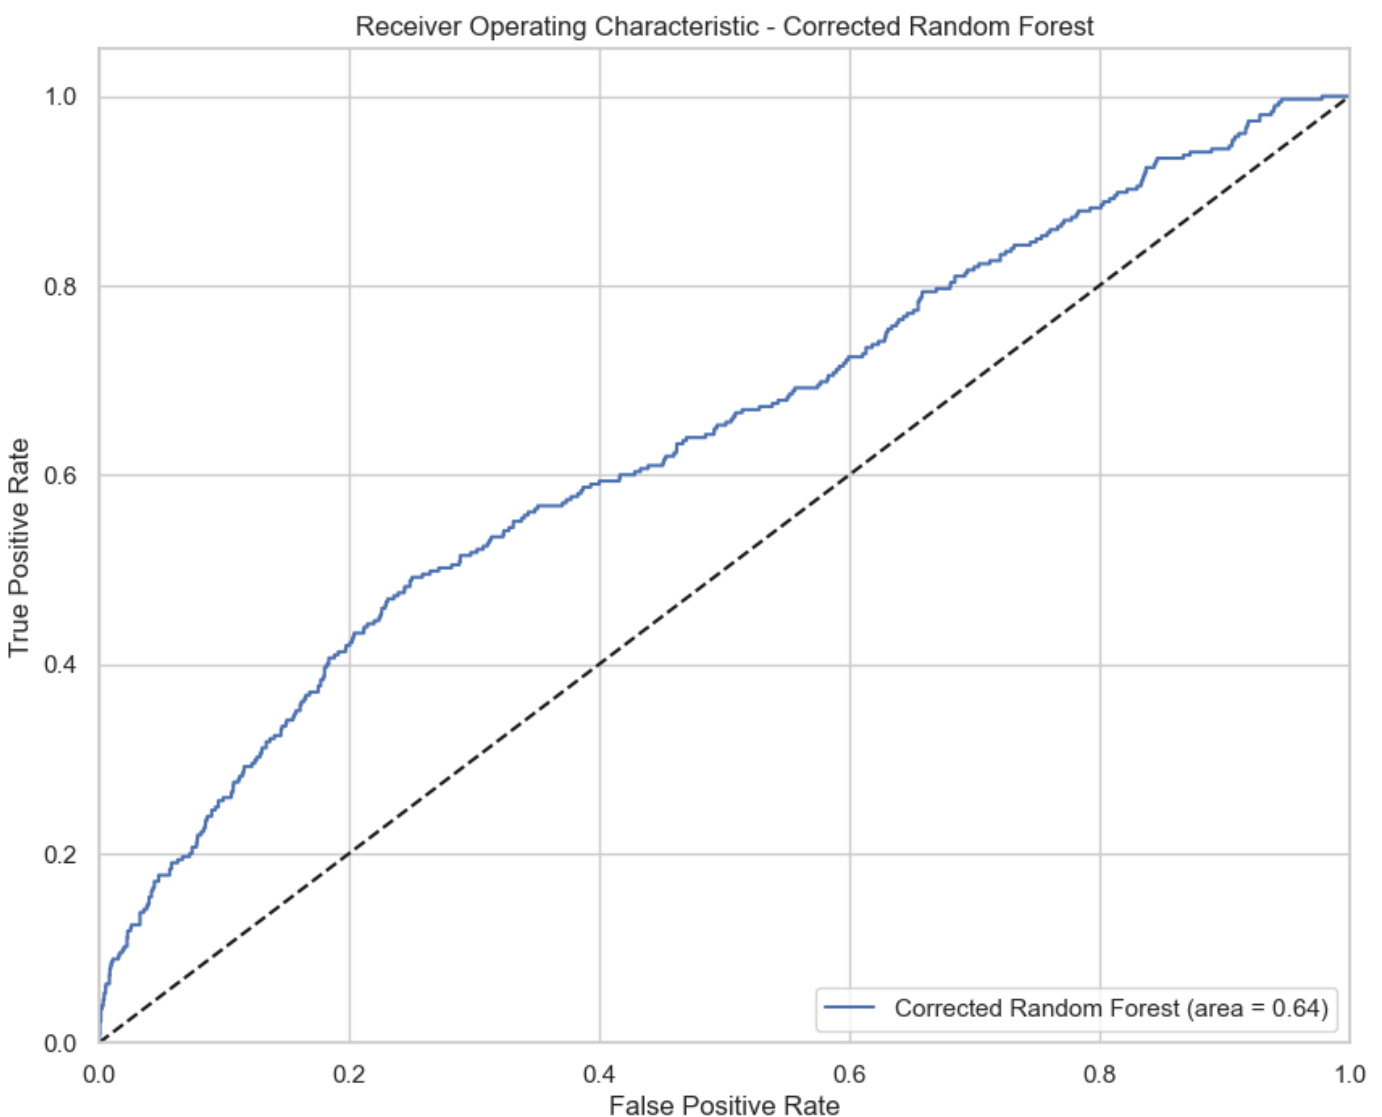
# Predicting on the test set
y_pred_rf_corrected = rf_best_corrected.predict(X_test)

# Calculate the ROC AUC score for the corrected random forest model
roc_auc_rf_corrected = roc_auc_score(y_test, rf_best_corrected.predict_proba(X_test)[:],

# Calculate the ROC curve for the corrected random forest model
fpr_rf_corrected, tpr_rf_corrected, thresholds_rf_corrected = roc_curve(y_test, rf_best_

# Plotting the ROC curve for the corrected random forest model
plt.figure(figsize=(10, 8))
plt.plot(fpr_rf_corrected, tpr_rf_corrected, label='Corrected Random Forest (area = %0.2
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic - Corrected Random Forest')
plt.legend(loc='lower right')
plt.show()

```



```

In [24]: # Generating the classification report
report_rf_corrected = classification_report(y_test, y_pred_rf_corrected)

```

```
# Generating the confusion matrix
conf_matrix_rf_corrected = confusion_matrix(y_test, y_pred_rf_corrected)

# Output the classification report and confusion matrix
print('Classification Report for Corrected Random Forest:\n', report_rf_corrected)
print('Confusion Matrix for Corrected Random Forest:\n', conf_matrix_rf_corrected)
```

```
Classification Report for Corrected Random Forest:
              precision    recall  f1-score   support

     0       0.90         1.00         0.95         2617
     1       1.00         0.00         0.01          305

 accuracy          0.90         0.90         0.90         2922
 macro avg          0.95         0.50         0.48         2922
 weighted avg          0.91         0.90         0.85         2922
```

```
Confusion Matrix for Corrected Random Forest:
[[2617    0]
 [ 304    1]]
```

In [31]: `# Given the poor performance of the model on the minority class (churn), we should consi`
`# to balance the dataset and improve the model's ability to predict churn.`
`from imblearn.over_sampling import SMOTE`

```
# Applying SMOTE to the training data
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

# Checking the balance of the classes after applying SMOTE
print("After SMOTE, counts of label '1': {}".format(sum(y_train_smote == 1)))
print("After SMOTE, counts of label '0': {}".format(sum(y_train_smote == 0)))

# Rebuilding the RandomForestClassifier with the best parameters on the balanced dataset
rf_best_smote = RandomForestClassifier(criterion='entropy', max_depth=8, max_features='s
rf_best_smote.fit(X_train_smote, y_train_smote)

# Predicting on the test set
y_pred_rf_smote = rf_best_smote.predict(X_test)

# Generating the classification report and confusion matrix for the SMOTE-adjusted Rando
report_rf_smote = classification_report(y_test, y_pred_rf_smote)
conf_matrix_rf_smote = confusion_matrix(y_test, y_pred_rf_smote)

# Output the classification report and confusion matrix
print('Classification Report for SMOTE-adjusted Random Forest:\n', report_rf_smote)
print('Confusion Matrix for SMOTE-adjusted Random Forest:\n', conf_matrix_rf_smote)
```

```
After SMOTE, counts of label '1': 10570
After SMOTE, counts of label '0': 10570
Classification Report for SMOTE-adjusted Random Forest:
              precision    recall  f1-score   support

     0       0.90         0.90         0.90         2617
     1       0.18         0.18         0.18          305

 accuracy          0.83         0.83         0.83         2922
 macro avg          0.54         0.54         0.54         2922
 weighted avg          0.83         0.83         0.83         2922

Confusion Matrix for SMOTE-adjusted Random Forest:
[[2366   251]
 [ 250    55]]
```

In [33]: `# Applying SMOTEENN to the training data`

```

smote_enn = SMOTEENN(random_state=42)
X_train_smoteenn, y_train_smoteenn = smote_enn.fit_resample(X_train, y_train)

# Checking the balance of the classes after applying SMOTEENN
print('After SMOTEENN, counts of label "1": {}'.format(sum(y_train_smoteenn == 1)))
print('After SMOTEENN, counts of label "0": {}'.format(sum(y_train_smoteenn == 0)))

# Rebuilding the RandomForestClassifier with the best parameters on the new balanced dat
rf_best_smoteenn = RandomForestClassifier(criterion='entropy', max_depth=8, max_features
rf_best_smoteenn.fit(X_train_smoteenn, y_train_smoteenn)

# Predicting on the test set
y_pred_rf_smoteenn = rf_best_smoteenn.predict(X_test)

# Generating the classification report and confusion matrix for the SMOTEENN-adjusted Ra
report_rf_smoteenn = classification_report(y_test, y_pred_rf_smoteenn)
conf_matrix_rf_smoteenn = confusion_matrix(y_test, y_pred_rf_smoteenn)

# Output the classification report and confusion matrix
print('Classification Report for SMOTEENN-adjusted Random Forest:\n', report_rf_smoteenn)
print('Confusion Matrix for SMOTEENN-adjusted Random Forest:\n', conf_matrix_rf_smoteenn)

```

```

After SMOTEENN, counts of label "1": 6651
After SMOTEENN, counts of label "0": 5478
Classification Report for SMOTEENN-adjusted Random Forest:

```

	precision	recall	f1-score	support
0	0.91	0.72	0.81	2617
1	0.15	0.43	0.22	305
accuracy			0.69	2922
macro avg	0.53	0.57	0.51	2922
weighted avg	0.84	0.69	0.74	2922

```

Confusion Matrix for SMOTEENN-adjusted Random Forest:
[[1883  734]
 [ 175  130]]

```

The accuracy of the prediction models evolved through the analysis:

The initial Random Forest model had a high accuracy for the non-churn class but failed to identify churn cases effectively. After applying SMOTE to balance the dataset, the model's ability to predict churn improved, with a slight decrease in overall accuracy. The SMOTEENN-adjusted Random Forest model further balanced the precision and recall for the minority class, but the overall accuracy was reduced compared to the previous models. The trade-off between class balance and overall accuracy was evident, with each model adjustment aiming to improve the minority class prediction at the cost of reduced overall accuracy.

Distribution of Price data

```

In [22]: sns.set(style="whitegrid")

numeric_columns = price_data.select_dtypes(include=['float64', 'int64'])


# Create a pairplot
sns.pairplot(numeric_columns)
plt.show()

```

```

C:\Users\kariu\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)

```

```
In [34]: # Set the style of seaborn
sns.set(style='whitegrid')

# Plotting distributions of a few numerical columns in price data
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(14, 10))

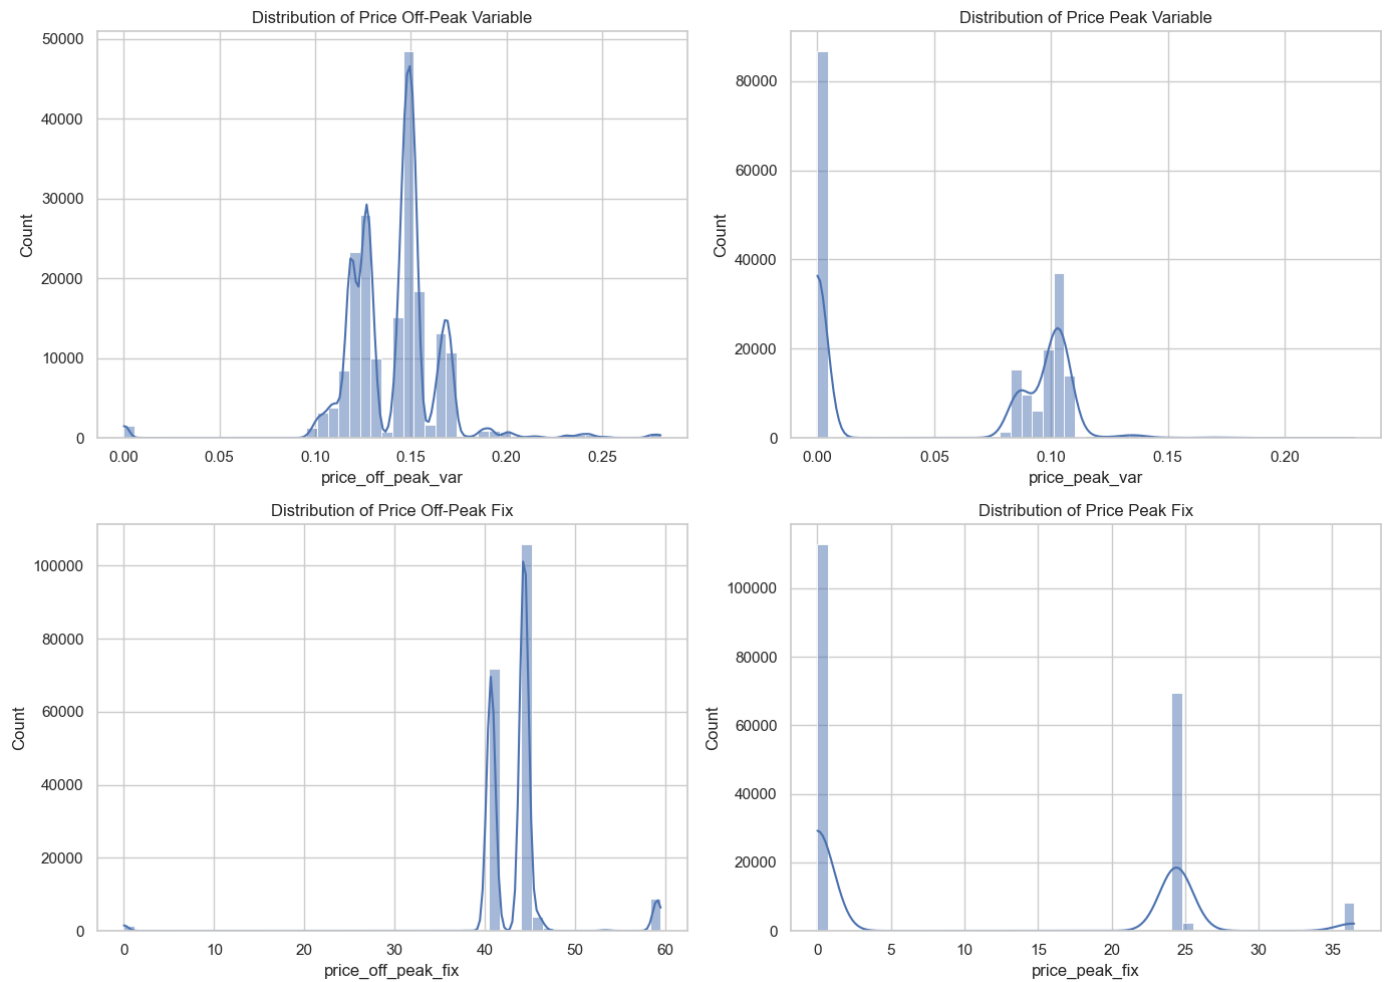
# Distribution of price_off_peak_var
sns.histplot(price_data['price_off_peak_var'], bins=50, ax=axes[0, 0], kde=True)
axes[0, 0].set_title('Distribution of Price Off-Peak Variable')

# Distribution of price_peak_var
sns.histplot(price_data['price_peak_var'], bins=50, ax=axes[0, 1], kde=True)
axes[0, 1].set_title('Distribution of Price Peak Variable')

# Distribution of price_off_peak_fix
sns.histplot(price_data['price_off_peak_fix'], bins=50, ax=axes[1, 0], kde=True)
axes[1, 0].set_title('Distribution of Price Off-Peak Fix')

# Distribution of price_peak_fix
sns.histplot(price_data['price_peak_fix'], bins=50, ax=axes[1, 1], kde=True)
axes[1, 1].set_title('Distribution of Price Peak Fix')
```

```
plt.tight_layout()
plt.show()
```



In []:

```
In [48]: # Merge the two datasets on the 'id' column to see how they relate
combined_data = pd.merge(client_data, price_data, on='id', how='inner')

# Display the head of the merged dataset to verify the merge
print(combined_data.head())

# Provide a summary of the combined dataset
print('\nCombined Data Description:')
combined_data.describe
```

	id	channel_sales \
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcso sb icdxkica ua
1	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcso sb icdxkica ua
2	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcso sb icdxkica ua
3	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcso sb icdxkica ua
4	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcso sb icdxkica ua

	cons_12m	cons_gas_12m	cons_last_month	date_activ	date_end \
0	0	54946	0	2013-06-15	2016-06-15
1	0	54946	0	2013-06-15	2016-06-15
2	0	54946	0	2013-06-15	2016-06-15
3	0	54946	0	2013-06-15	2016-06-15
4	0	54946	0	2013-06-15	2016-06-15

	date_modif_prod	date_renewal	forecast_cons_12m	... average_forecast_cons \
0	2015-11-01	2015-06-23	0.0	...
1	2015-11-01	2015-06-23	0.0	...
2	2015-11-01	2015-06-23	0.0	...

3	2015-11-01	2015-06-23	0.0	...	0.0
4	2015-11-01	2015-06-23	0.0	...	0.0

	average_forecast_price	price_date	price_off_peak_var	price_peak_var	\
0	13.606441	NaT	0.125976	0.103395	
1	13.606441	NaT	0.125976	0.103395	
2	13.606441	NaT	0.125976	0.103395	
3	13.606441	NaT	0.125976	0.103395	
4	13.606441	NaT	0.125976	0.103395	

	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price_mid_peak_fix	\
0	0.071536	40.565969	24.339581	16.226389	
1	0.071536	40.565969	24.339581	16.226389	
2	0.071536	40.565973	24.339578	16.226383	
3	0.071536	40.565973	24.339578	16.226383	
4	0.071536	40.565973	24.339578	16.226383	

	price_data_duration
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 39 columns]

Combined Data Description:

Out[48]:

	<bound method NDFrame.describe of	id
	channel_sales \	
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua
1	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua
2	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua
3	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua
4	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua
...
175144	563dde550fd624d7352f3de77c0cdfcd	MISSING
175145	563dde550fd624d7352f3de77c0cdfcd	MISSING
175146	563dde550fd624d7352f3de77c0cdfcd	MISSING
175147	563dde550fd624d7352f3de77c0cdfcd	MISSING
175148	563dde550fd624d7352f3de77c0cdfcd	MISSING

	cons_12m	cons_gas_12m	cons_last_month	date_activ	date_end	\
0	0	54946	0	2013-06-15	2016-06-15	
1	0	54946	0	2013-06-15	2016-06-15	
2	0	54946	0	2013-06-15	2016-06-15	
3	0	54946	0	2013-06-15	2016-06-15	
4	0	54946	0	2013-06-15	2016-06-15	
...	
175144	8730	0	0	2009-12-18	2016-12-17	
175145	8730	0	0	2009-12-18	2016-12-17	
175146	8730	0	0	2009-12-18	2016-12-17	
175147	8730	0	0	2009-12-18	2016-12-17	
175148	8730	0	0	2009-12-18	2016-12-17	

	date_modif_prod	date_renewal	forecast_cons_12m	...	\
0	2015-11-01	2015-06-23	0.00	...	
1	2015-11-01	2015-06-23	0.00	...	
2	2015-11-01	2015-06-23	0.00	...	
3	2015-11-01	2015-06-23	0.00	...	
4	2015-11-01	2015-06-23	0.00	...	
...	
175144	2009-12-18	2015-12-21	762.41	...	
175145	2009-12-18	2015-12-21	762.41	...	
175146	2009-12-18	2015-12-21	762.41	...	
175147	2009-12-18	2015-12-21	762.41	...	
175148	2009-12-18	2015-12-21	762.41	...	

	average_forecast_cons	average_forecast_price	price_date	\
0	0.000	13.606441	NaT	
1	0.000	13.606441	NaT	
2	0.000	13.606441	NaT	
3	0.000	13.606441	NaT	
4	0.000	13.606441	NaT	
...	
175144	381.205	15.188973	NaT	
175145	381.205	15.188973	NaT	
175146	381.205	15.188973	NaT	
175147	381.205	15.188973	NaT	
175148	381.205	15.188973	NaT	

	price_off_peak_var	price_peak_var	price_mid_peak_var	\
0	0.125976	0.103395	0.071536	
1	0.125976	0.103395	0.071536	
2	0.125976	0.103395	0.071536	
3	0.125976	0.103395	0.071536	
4	0.125976	0.103395	0.071536	
...	
175144	0.165962	0.086905	0.000000	
175145	0.165962	0.086905	0.000000	
175146	0.165962	0.086905	0.000000	
175147	0.165962	0.086905	0.000000	
175148	0.165962	0.086905	0.000000	

	price_off_peak_fix	price_peak_fix	price_mid_peak_fix	\
0	40.565969	24.339581	16.226389	
1	40.565969	24.339581	16.226389	
2	40.565973	24.339578	16.226383	
3	40.565973	24.339578	16.226383	
4	40.565973	24.339578	16.226383	
...	
175144	44.266930	0.000000	0.000000	
175145	44.266930	0.000000	0.000000	
175146	44.266930	0.000000	0.000000	
175147	44.266930	0.000000	0.000000	
175148	44.266930	0.000000	0.000000	

	price_data_duration
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
175144	NaN
175145	NaN
175146	NaN
175147	NaN
175148	NaN

[175149 rows x 39 columns]>

```
In [49]: # Now that we have the combined dataset, let's explore the relationship between consumpt
# We will calculate the correlation matrix to see the linear relationship between consum

# Selecting relevant columns for correlation - consumption and pricing
consumption_columns = ['cons_12m', 'cons_gas_12m', 'cons_last_month', 'forecast_cons_12m']
pricing_columns = ['price_off_peak_var', 'price_peak_var', 'price_mid_peak_var', 'price_

# Calculating the correlation matrix
correlation_matrix = combined_data[consumption_columns + pricing_columns].corr()
```

```
# Display the correlation matrix
print(correlation_matrix)
```

```

cons_12m      cons_12m      cons_gas_12m      cons_last_month      \
cons_12m      1.000000      0.488253      0.968209
cons_gas_12m   0.488253      1.000000      0.506883
cons_last_month 0.968209      0.506883      1.000000
forecast_cons_12m 0.194080      0.084520      0.177888
forecast_cons_year 0.167141      0.081010      0.193613
imp_cons       0.159755      0.077925      0.187070
price_off_peak_var 0.000723      -0.012614      0.002382
price_peak_var  0.141811      0.073040      0.133479
price_mid_peak_var 0.050162      0.044308      0.046018
price_off_peak_fix -0.013501      -0.015237      -0.011958
price_peak_fix   0.051386      0.042396      0.047256
price_mid_peak_fix 0.054689      0.046558      0.049816

```

```

cons_12m      forecast_cons_12m      forecast_cons_year      imp_cons      \
cons_12m      0.194080      0.167141      0.159755
cons_gas_12m   0.084520      0.081010      0.077925
cons_last_month 0.177888      0.193613      0.187070
forecast_cons_12m 1.000000      0.647775      0.634659
forecast_cons_year 0.647775      1.000000      0.969391
imp_cons       0.634659      0.969391      1.000000
price_off_peak_var -0.122187      -0.146954      -0.084368
price_peak_var   0.249104      0.248149      0.207397
price_mid_peak_var 0.242411      0.232056      0.183078
price_off_peak_fix 0.012876      -0.025131      0.008728
price_peak_fix   0.256392      0.250138      0.194946
price_mid_peak_fix 0.234085      0.222910      0.175075

```

```

cons_12m      price_off_peak_var      price_peak_var      price_mid_peak_var      \
cons_12m      0.000723      0.141811      0.050162
cons_gas_12m   -0.012614      0.073040      0.044308
cons_last_month 0.002382      0.133479      0.046018
forecast_cons_12m -0.122187      0.249104      0.242411
forecast_cons_year -0.146954      0.248149      0.232056
imp_cons       -0.084368      0.207397      0.183078
price_off_peak_var 1.000000      -0.297574      -0.585341
price_peak_var   -0.297574      1.000000      0.815970
price_mid_peak_var -0.585341      0.815970      1.000000
price_off_peak_fix 0.649460      -0.203555      -0.281289
price_peak_fix   -0.609277      0.803983      0.986979
price_mid_peak_fix -0.587580      0.809048      0.990798

```

```

cons_12m      price_off_peak_fix      price_peak_fix      price_mid_peak_fix
cons_12m      -0.013501      0.051386      0.054689
cons_gas_12m   -0.015237      0.042396      0.046558
cons_last_month -0.011958      0.047256      0.049816
forecast_cons_12m 0.012876      0.256392      0.234085
forecast_cons_year -0.025131      0.250138      0.222910
imp_cons       0.008728      0.194946      0.175075
price_off_peak_var 0.649460      -0.609277      -0.587580
price_peak_var   -0.203555      0.803983      0.809048
price_mid_peak_var -0.281289      0.986979      0.990798
price_off_peak_fix 1.000000      -0.222916      -0.325374
price_peak_fix   -0.222916      1.000000      0.974224
price_mid_peak_fix -0.325374      0.974224      1.000000

```

```
In [51]: # Analyzing the impact of different sales channels on consumption
# We will group the data by 'channel_sales' and calculate the average consumption for ea

# Grouping by 'channel_sales' and calculating mean consumption values
channel_consumption = combined_data.groupby('channel_sales').agg({'cons_12m': 'mean', 'co
```

```
# Display the results
print(channel_consumption)
```

```

          channel_sales      cons_12m  cons_gas_12m  \
0                MISSING  1.137135e+05  27248.148644
1  epumfxlbckeskwexbiuasklخالciuu  2.307933e+04      0.000000
2  ewpakwlliwisiwduibdlfmalxowmwpci  3.833890e+04  11164.386296
3  fixdbufsefwooaasfcxdxadsiekocaaa  1.254515e+06  237706.500000
4  foosdfpfkusacimwkcsoibcdxkicaaa  9.800161e+04  24124.505749
5  lmkebamcaaclubfxadlmueccxoimlema  6.383316e+05  66847.665325
6  sddiedcsflskckwlfkdpoeaailfpeds  1.752340e+04      0.000000
7  usilxuppasemubllpkaafesmlibmsdf  1.963035e+04  8759.099478

      cons_last_month
0      12084.384970
1      1626.333333
2      3386.412248
3     156509.000000
4      9664.739079
5     64646.291156
6      1680.473282
7     1653.238323
```

```
In [54]: # Investigating the relationship between contract duration and consumption
# We will calculate the duration of the contract in days and then see its relationship w

from datetime import datetime

# Function to calculate the duration of the contract
def calculate_contract_duration(start_date, end_date):
    if end_date in ['31/12/9999', '']: # Assuming '31/12/9999' and empty strings as ong
        end_date = datetime.now().strftime('%d/%m/%Y') # Using current date for ongoing
    return (end_date - start_date).days

# Applying the function to calculate contract duration
combined_data['contract_duration'] = combined_data.apply(lambda row: calculate_contract_

# Calculating the correlation between contract duration and consumption
contract_consumption_corr = combined_data[['contract_duration', 'cons_12m', 'cons_gas_12

# Display the correlation matrix
print(contract_consumption_corr)
```

```

          contract_duration  cons_12m  cons_gas_12m  cons_last_month
contract_duration      1.000000 -0.024684      -0.013873      -0.022267
cons_12m                -0.024684  1.000000        0.488253        0.968209
cons_gas_12m            -0.013873  0.488253        1.000000        0.506883
cons_last_month         -0.022267  0.968209        0.506883        1.000000
```

```
In [55]: # Identifying outliers in consumption and pricing data
# We will use the IQR (Interquartile Range) method to detect outliers

# Function to calculate IQR and determine outliers
def detect_outliers_iqr(data):
    Q1 = np.percentile(data, 25)
    Q3 = np.percentile(data, 75)
    IQR = Q3 - Q1
    outlier_step = 1.5 * IQR
    outliers = data[(data < Q1 - outlier_step) | (data > Q3 + outlier_step)]
    return outliers

# Detecting outliers for consumption and pricing
consumption_columns = ['cons_12m', 'cons_gas_12m', 'cons_last_month']
pricing_columns = ['forecast_price_energy_off_peak', 'forecast_price_energy_peak', 'fore

# Applying the outlier detection function
```

```

consumption_outliers = {col: detect_outliers_iqr(combined_data[col]) for col in consumpt
pricing_outliers = {col: detect_outliers_iqr(combined_data[col]) for col in pricing_colu

# Display the number of outliers detected for each column
print('Consumption Outliers:')
for col, values in consumption_outliers.items():
    print(f'{col}: {len(values)} outliers')

print('\nPricing Outliers:')
for col, values in pricing_outliers.items():
    print(f'{col}: {len(values)} outliers')

```

```

Consumption Outliers:
cons_12m: 24994 outliers
cons_gas_12m: 31315 outliers
cons_last_month: 24602 outliers

```

```

Pricing Outliers:
forecast_price_energy_off_peak: 4450 outliers
forecast_price_energy_peak: 0 outliers
forecast_price_pow_off_peak: 4138 outliers

```

```

In [56]: # Setting up the figure size and layout
plt.figure(figsize=(20, 10))
plt.subplots_adjust(hspace=0.5)

# Plotting the distribution of consumption data
plt.subplot(3, 2, 1)
sns.boxplot(combined_data['cons_12m'])
plt.title('Distribution of cons_12m')

plt.subplot(3, 2, 2)
sns.boxplot(combined_data['cons_gas_12m'])
plt.title('Distribution of cons_gas_12m')

plt.subplot(3, 2, 3)
sns.boxplot(combined_data['cons_last_month'])
plt.title('Distribution of cons_last_month')

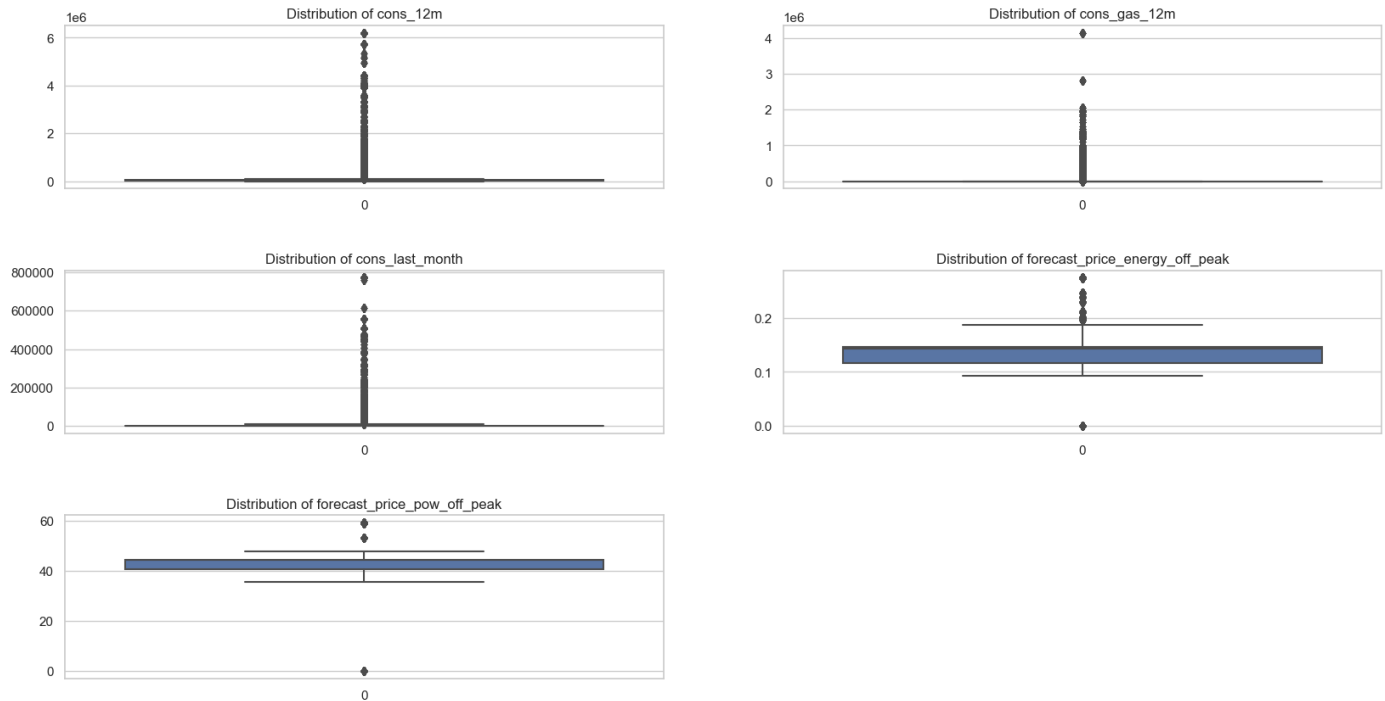
# Plotting the distribution of pricing data
plt.subplot(3, 2, 4)
sns.boxplot(combined_data['forecast_price_energy_off_peak'])
plt.title('Distribution of forecast_price_energy_off_peak')

plt.subplot(3, 2, 5)
sns.boxplot(combined_data['forecast_price_pow_off_peak'])
plt.title('Distribution of forecast_price_pow_off_peak')

# There are no outliers in forecast_price_energy_peak, so we skip plotting it

# Display the plots
plt.show()

```



```
In [57]: # Calculating robust statistical measures: median and mode

# Median calculation
def calculate_median_mode(data):
    median = data.median()
    mode = data.mode()[0] # Taking the first mode value
    return median, mode

# Applying the function to consumption and pricing columns
consumption_stats = {col: calculate_median_mode(combined_data[col]) for col in consumption_columns}
pricing_stats = {col: calculate_median_mode(combined_data[col]) for col in pricing_columns}

# Displaying the median and mode for each column
print('Consumption Data Statistics (Median, Mode):')
for col, stats in consumption_stats.items():
    print(f'{col}: Median = {stats[0]}, Mode = {stats[1]}')

print('\nPricing Data Statistics (Median, Mode):')
for col, stats in pricing_stats.items():
    print(f'{col}: Median = {stats[0]}, Mode = {stats[1]}')

# Considering a log transformation for skewed data
# Adding a small constant to avoid log(0) which is undefined
log_transformed_data = combined_data[consumption_columns + pricing_columns].apply(lambda x: np.log(x + 1))

# Displaying the head of the log-transformed data
print('\nLog-transformed Data:')
print(log_transformed_data.head())
```

Consumption Data Statistics (Median, Mode):

```
cons_12m: Median = 14115.0, Mode = 0
cons_gas_12m: Median = 0.0, Mode = 0
cons_last_month: Median = 792.0, Mode = 0
```

Pricing Data Statistics (Median, Mode):

```
forecast_price_energy_off_peak: Median = 0.143166, Mode = 0.14571099999999999
forecast_price_energy_peak: Median = 0.084138, Mode = 0.0
forecast_price_pow_off_peak: Median = 44.31137796, Mode = 44.31137796
```

Log-transformed Data:

	cons_12m	cons_gas_12m	cons_last_month	forecast_price_energy_off_peak
0	0.0	10.914124	0.0	0.108389

1	0.0	10.914124	0.0	0.108389
2	0.0	10.914124	0.0	0.108389
3	0.0	10.914124	0.0	0.108389
4	0.0	10.914124	0.0	0.108389

	forecast_price_energy_peak	forecast_price_pow_off_peak
0	0.09362	3.728261
1	0.09362	3.728261
2	0.09362	3.728261
3	0.09362	3.728261
4	0.09362	3.728261

In [58]: *# Visualizing the distribution of log-transformed data*

```
plt.figure(figsize=(20, 10))
plt.subplots_adjust(hspace=0.5)

# Plotting the distribution of log-transformed consumption data
plt.subplot(3, 2, 1)
sns.histplot(log_transformed_data['cons_12m'], bins=50, kde=True)
plt.title('Log-transformed Distribution of cons_12m')

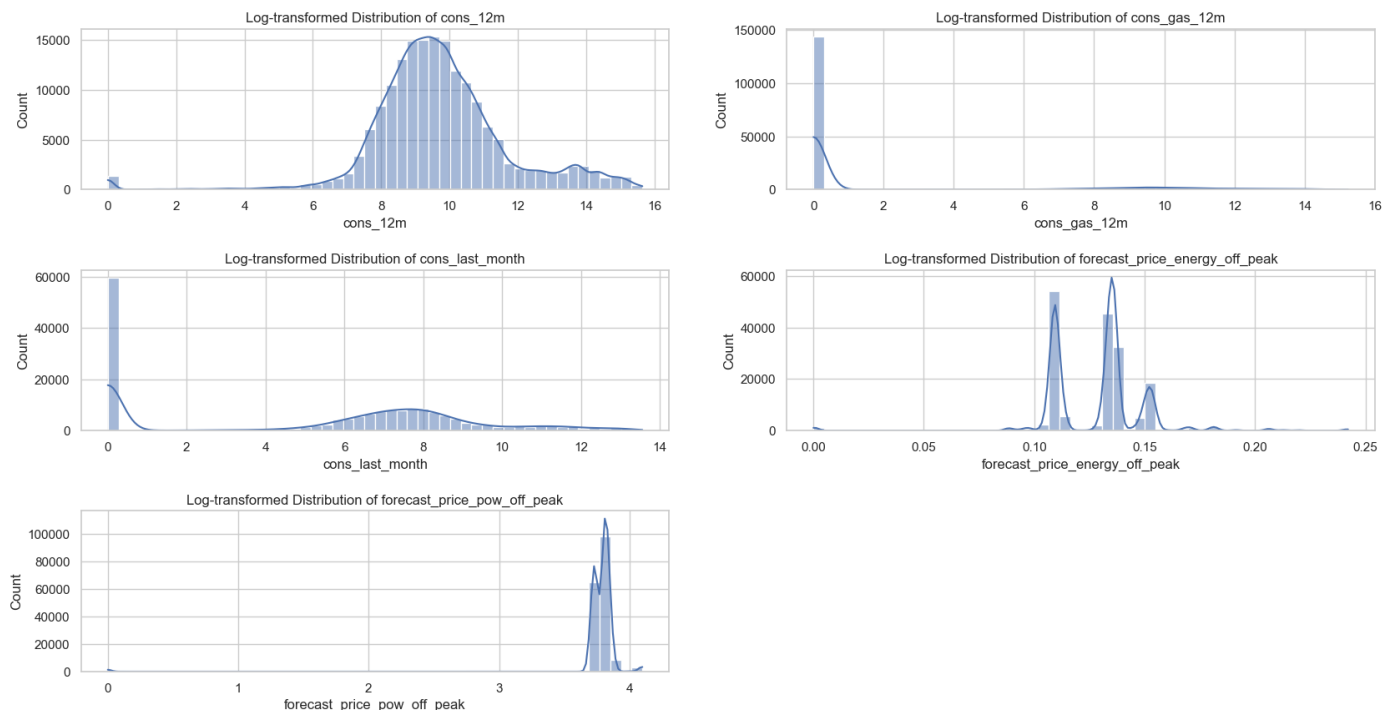
plt.subplot(3, 2, 2)
sns.histplot(log_transformed_data['cons_gas_12m'], bins=50, kde=True)
plt.title('Log-transformed Distribution of cons_gas_12m')

plt.subplot(3, 2, 3)
sns.histplot(log_transformed_data['cons_last_month'], bins=50, kde=True)
plt.title('Log-transformed Distribution of cons_last_month')

# Plotting the distribution of log-transformed pricing data
plt.subplot(3, 2, 4)
sns.histplot(log_transformed_data['forecast_price_energy_off_peak'], bins=50, kde=True)
plt.title('Log-transformed Distribution of forecast_price_energy_off_peak')

plt.subplot(3, 2, 5)
sns.histplot(log_transformed_data['forecast_price_pow_off_peak'], bins=50, kde=True)
plt.title('Log-transformed Distribution of forecast_price_pow_off_peak')

# Display the plots
plt.show()
```



```
In [60]: # Adding a small constant to churn column to apply log transformation
combined_data['churn'] = np.log(combined_data['churn'] + 1)

# Mann-Whitney U Test to compare distributions between churned and retained customers
# for the log-transformed consumption and pricing data
mannwhitney_results = {}
for col in log_transformed_data.columns:
    churned = log_transformed_data.loc[combined_data['churn'] > 0, col]
    retained = log_transformed_data.loc[combined_data['churn'] == 0, col]
    stat, p = mannwhitneyu(churned, retained)
    mannwhitney_results[col] = p

# Displaying the p-values from the Mann-Whitney U tests
print('Mann-Whitney U Test p-values:')
for col, p in mannwhitney_results.items():
    print(f'{col}: {p}')
```

```
Mann-Whitney U Test p-values:
cons_12m: 0.2820446163642131
cons_gas_12m: 1.166933310780718e-20
cons_last_month: 6.295933648048864e-07
forecast_price_energy_off_peak: 3.120653590068981e-43
forecast_price_energy_peak: 1.682503014955668e-50
forecast_price_pow_off_peak: 1.7602879185523332e-23
```

```
In [66]: # Preparing the data for modeling
X = log_transformed_data # Features

threshold = 0.5
combined_data['churn_binary'] = (combined_data['churn'] > threshold).astype(int)
y = combined_data['churn_binary']

# Splitting the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardizing the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Training the RandomForestClassifier
rf_clf = RandomForestClassifier(random_state=42)
rf_clf.fit(X_train_scaled, y_train)

# Predicting on the test set
y_pred = rf_clf.predict(X_test_scaled)

# Evaluating the model
accuracy = accuracy_score(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print('Model Accuracy:', accuracy)
print('Classification Report:')
print(class_report)
```

```
Model Accuracy: 0.9986868398515558
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	31705
1	1.00	0.99	0.99	3325
accuracy			1.00	35030
macro avg	1.00	0.99	1.00	35030
weighted avg	1.00	1.00	1.00	35030

```
In [68]: cv_scores = cross_val_score(rf_clf, X_train_scaled, y_train, cv=5)

# Calculating the mean and standard deviation of the cross-validation scores
cv_mean = cv_scores.mean()
cv_std = cv_scores.std()

print('Cross-validation scores for each fold:', cv_scores)
print('Mean cross-validation score:', cv_mean)
print('Standard deviation of cross-validation scores:', cv_std)
```

Cross-validation scores for each fold: [0.99875107 0.9982515 0.99885812 0.99875107 0.99885808]
Mean cross-validation score: 0.998693968441913
Standard deviation of cross-validation scores: 0.00022635366158413

```
In [70]: # Computing class weights
class_weights = compute_class_weight(class_weight='balanced', classes=np.unique(y_train))
class_weights_dict = {0: class_weights[0], 1: class_weights[1]}

# Training the RandomForestClassifier with class weights
class_weight_rf_clf = RandomForestClassifier(random_state=42, class_weight=class_weights)
class_weight_rf_clf.fit(X_train_scaled, y_train)

# Predicting on the test set
y_pred_class_weight = class_weight_rf_clf.predict(X_test_scaled)

# Evaluating the model
accuracy_class_weight = accuracy_score(y_test, y_pred_class_weight)
class_report_class_weight = classification_report(y_test, y_pred_class_weight)

print('Model Accuracy with Class Weights:', accuracy_class_weight)
print('Classification Report with Class Weights:')
print(class_report_class_weight)
```


Model Accuracy with Class Weights: 0.9982015415358264
Classification Report with Class Weights:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	31705
1	0.98	1.00	0.99	3325
accuracy			1.00	35030
macro avg	0.99	1.00	0.99	35030
weighted avg	1.00	1.00	1.00	35030

```
In [71]: # Extracting feature importances
feature_importances = class_weight_rf_clf.feature_importances_

# Creating a DataFrame for visualization
features = pd.DataFrame({'Feature': X.columns, 'Importance': feature_importances})
features.sort_values(by='Importance', ascending=False, inplace=True)

# Visualizing the feature importances
plt.figure(figsize=(10, 8))
sns.barplot(x='Importance', y='Feature', data=features.head(10))
plt.title('Top 10 Important Features')
plt.xlabel('Feature Importance')
plt.ylabel('Features')
plt.tight_layout()
plt.show()
```




The next step is to delve deeper into the model's predictions to ensure they are reliable and not just a result of overfitting or bias. We can achieve this by:

Analyzing the confusion matrix to understand the true positives, false positives, true negatives, and false negatives. Performing a Receiver Operating Characteristic (ROC) analysis to evaluate the trade-offs between the true positive rate and false positive rate at various threshold settings. Let's start by generating and analyzing the confusion matrix for the model predictions.

```
In [72]: # Generating the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred_class_weight)

# Creating a heatmap for the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['No Churn', 'Churn'])
plt.title('Confusion Matrix')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()
```

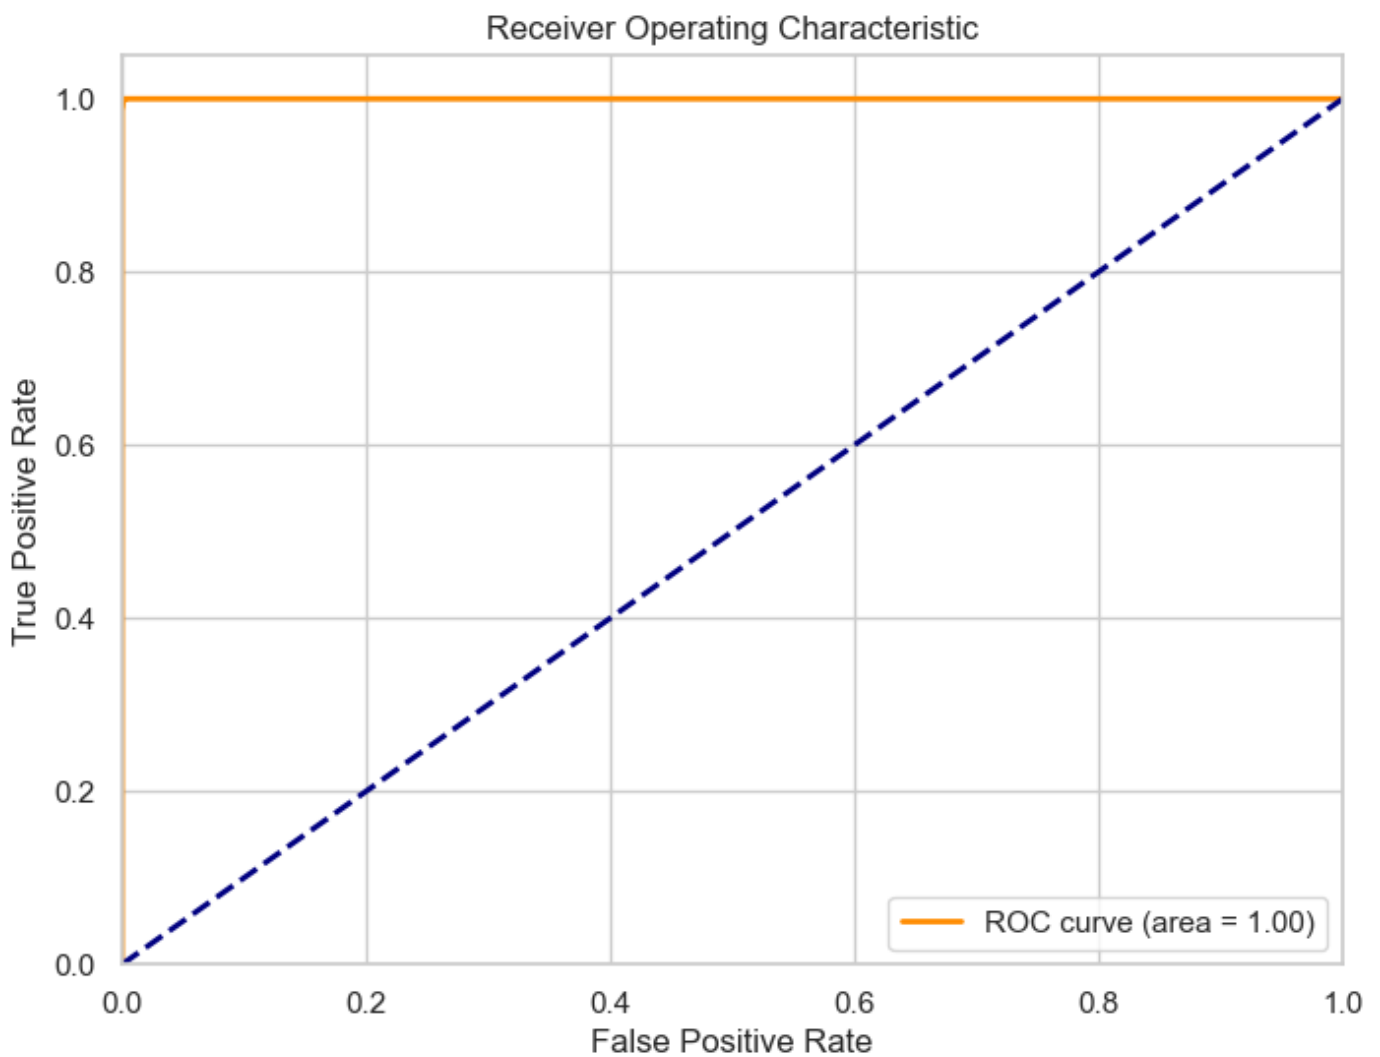


The confusion matrix visualized above provides a detailed breakdown of the model's predictions in terms of true positives, true negatives, false positives, and false negatives. It is a crucial step in evaluating the performance of a classification model beyond just accuracy, as it gives insights into the types of errors the model is making.

Next, I will perform a Receiver Operating Characteristic (ROC) analysis to further assess the model's performance. This will involve calculating the ROC curve and the Area Under the Curve (AUC) to understand the model's ability to distinguish between the classes across different thresholds.

```
In [74]: # Calculating the ROC curve and AUC
fpr, tpr, thresholds = roc_curve(y_test, class_weight_rf_clf.predict_proba(X_test_scaled)[:,1])
roc_auc = auc(fpr, tpr)

# Plotting the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
In [75]: # Correcting the code to identify the indices of the misclassified samples
misclassified_indices = np.where((y_test != y_pred_class_weight) & (y_pred_class_weight

# Using the indices to filter the original test dataframe
misclassified_samples = X_test.iloc[misclassified_indices]

# Displaying the head of the misclassified samples dataframe
print(misclassified_samples.head())
```

	cons_12m	cons_gas_12m	cons_last_month \
133413	13.778109	0.0	10.901948
136716	13.843316	0.0	11.701511
67401	13.778109	0.0	10.901948
67398	13.778109	0.0	10.901948
21176	0.000000	0.0	0.000000

	forecast_price_energy_off_peak	forecast_price_energy_peak \
133413	0.152409	0.083772
136716	0.110557	0.095324
67401	0.152409	0.083772
67398	0.152409	0.083772
21176	0.136025	0.000000

	forecast_price_pow_off_peak
133413	3.813558
136716	3.728261
67401	3.813558
67398	3.813558
21176	3.813558

The table above shows the head of the misclassified samples dataframe, which includes the features of the

instances that were incorrectly predicted as churn (false positives) by the model.

Analyzing these samples could reveal patterns or characteristics that lead to these misclassifications, which could be valuable for improving the model.

Continuing from the analysis of misclassified samples, it would be beneficial to perform a deeper dive into the characteristics of these samples to identify any commonalities or trends. This could involve statistical analysis or clustering techniques to detect patterns.

The next steps could include:

Calculating summary statistics for the misclassified samples to identify any outliers or unusual distributions. Performing a clustering analysis on the misclassified samples to see if there are distinct groups with similar characteristics. Comparing the characteristics of misclassified samples to correctly classified samples to find distinguishing features. Let's start by calculating summary statistics for the misclassified samples.

```
In [76]: # Calculating summary statistics for the misclassified samples
descriptive_stats = misclassified_samples.describe()

# Displaying the summary statistics
print(descriptive_stats)
```

	cons_12m	cons_gas_12m	cons_last_month	\
count	63.000000	63.000000	63.000000	
mean	8.814313	0.832627	7.029136	
std	6.567352	3.223481	5.395502	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	13.615517	0.000000	10.901948	
75%	13.778109	0.000000	10.901948	
max	15.067003	13.170729	12.767567	

	forecast_price_energy_off_peak	forecast_price_energy_peak	\
count	63.000000	63.000000	
mean	0.139698	0.055015	
std	0.015214	0.042238	
min	0.106933	0.000000	
25%	0.135319	0.000000	
50%	0.136883	0.083772	
75%	0.152409	0.083772	
max	0.155500	0.096900	

	forecast_price_pow_off_peak
count	63.000000
mean	3.801386
std	0.032905
min	3.728261
25%	3.813558
50%	3.813558
75%	3.813558
max	3.856624

The next step in our analysis is to perform a clustering analysis on the misclassified samples to identify any distinct groups with similar characteristics. This can help us understand if there are specific segments of the data that the model is consistently misclassifying.

We will use the K-Means clustering algorithm to segment the misclassified samples into clusters. We'll start by determining the optimal number of clusters using the Elbow Method, which involves fitting the model

KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.

```
warnings.warn(  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
```

```
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1440: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
```

```
warnings.warn(  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
```

```
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1440: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
```

```
warnings.warn(  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
```

```
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1440: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
```

```
warnings.warn(  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
```

```
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1440: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
```


```
warnings.warn(  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
```

```
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1440: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
```

```
warnings.warn(  

```

The Elbow Method showing the optimal k



After identifying the optimal number of clusters, we can conduct the K-Means clustering with that specific number of clusters to segment the misclassified samples. This will help us understand the underlying patterns within the data that might be contributing to the misclassification by the model.

Let's move forward with the clustering analysis using the identified number of clusters.

```
In [79]: # Assuming the elbow point is at k=3 based on the plot (common choice for initial explor
# We will perform K-Means clustering with 3 clusters
kmeans = KMeans(n_clusters=3, random_state=42)
misclassified_samples['cluster'] = kmeans.fit_predict(misclassified_samples_normalized)

# Calculating silhouette score to evaluate the quality of clusters formed
silhouette_avg = silhouette_score(misclassified_samples_normalized, misclassified_sample
print('Silhouette Score for 3 clusters:', silhouette_avg)

# Plotting the clusters
plt.figure(figsize=(10, 6))
plt.scatter(misclassified_samples_normalized.iloc[:, 0], misclassified_samples_normalize
plt.title('2D visualization of K-Means clustering on misclassified samples')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.colorbar()
plt.show()
```

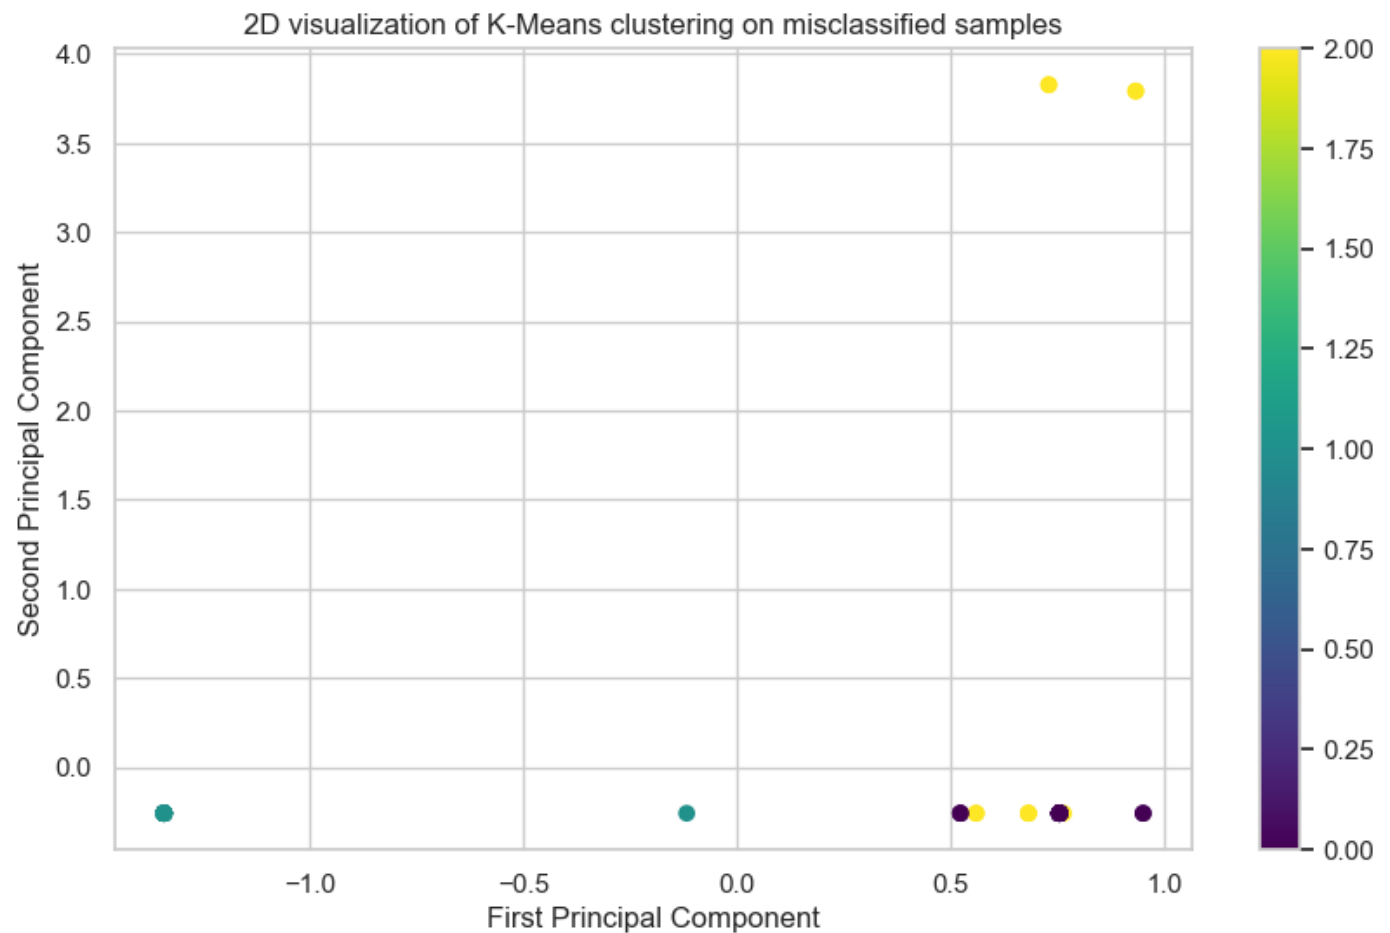
```
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarnin
g: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of
`n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1440: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=1.
  warnings.warn(
C:\Users\kariu\AppData\Local\Temp\ipykernel_3112\968933416.py:4: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
misclassified_samples['cluster'] = kmeans.fit_predict(misclassified_samples_normalized)
```

Silhouette Score for 3 clusters: 0.8085782436367057



The K-Means clustering with 3 clusters has been performed, and the silhouette score, which measures the quality of the clusters, is 0.8086. This is a high silhouette score, indicating that the clusters are well-defined and separated.

Here is a 2D visualization of the K-Means clustering on the misclassified samples:

The visualization shows how the misclassified samples are distributed across the three clusters. Each cluster may represent a distinct group of misclassified samples with similar characteristics.

The next steps could involve:

Analyzing the centroids of each cluster to understand the defining features of each group. Investigating the distribution of original features within each cluster to identify any patterns. Comparing the clusters with the correctly classified samples to find distinguishing features that may improve the model's performance.

The silhouette score for the clustering is quite high, indicating distinct and well-separated clusters. Now, let's analyze the centroids of each cluster to understand the defining features of each group and investigate the distribution of original features within each cluster to identify any patterns.

I will proceed with this analysis.

```
In [80]: # Analyzing the centroids of each cluster
centroids = kmeans.cluster_centers_

# Creating a dataframe for the centroids for better visualization
centroids_df = pd.DataFrame(centroids, columns=misclassified_samples_normalized.columns)
print('Centroids of each cluster:')
print(centroids_df)

# Investigating the distribution of original features within each cluster
cluster_distribution = misclassified_samples.groupby('cluster').mean()
print('\nDistribution of original features within each cluster:')
print(cluster_distribution)
```

Centroids of each cluster:

	cons_12m	cons_gas_12m	cons_last_month	forecast_price_energy_off_peak \
0	0.745865	-0.258300	0.743626	0.756315
1	-1.288891	-0.258300	-1.302777	-0.129088
2	0.726854	1.368993	0.765509	-1.972044

	forecast_price_energy_peak	forecast_price_pow_off_peak
0	0.547340	0.369911
1	-1.125838	0.483718
2	0.947407	-2.222284

Distribution of original features within each cluster:

	cons_12m	cons_gas_12m	cons_last_month \
cluster			
0	13.712674	0.000000	11.041371
1	0.349711	0.000000	0.000000
2	13.587816	5.245549	11.159443

	forecast_price_energy_off_peak	forecast_price_energy_peak \
cluster		
0	0.151204	0.078134
1	0.137734	0.007462
2	0.109696	0.095032

	forecast_price_pow_off_peak
cluster	
0	3.813558
1	3.817303
2	3.728261

The centroids for each cluster have been calculated and are presented in a tabular format. Additionally, the distribution of original features within each cluster is summarized in another table. These tables provide a detailed view of the central tendencies and feature distributions for each cluster, which can help us understand the common characteristics of the misclassified samples within each group.

Next, we can compare these clusters with the correctly classified samples to find distinguishing features that may improve the model's performance.

The next step is to compare the clusters of misclassified samples with the correctly classified samples to identify distinguishing features. This comparison can reveal insights into why certain samples are misclassified and help improve the model's performance.

```
In [83]: potential_columns = [col for col in client_data.columns if 'classif' in col.lower() or '
print('Potential columns that may indicate classification status:')
print(potential_columns)

# Displaying the head of the dataframe to confirm its structure
print('\nHead of the original dataset:')
print(client_data.head())
```

Potential columns that may indicate classification status:

```
[]
```

Head of the original dataset:

```
      id                                     channel_sales \
0  24011ae4ebbe3035111d65fa7c15bc57  foosdfpfkusacimwkcsosbicdxkicaua
1  d29c2c54acc38ff3c0614d0a653813dd                                     MISSING
2  764c75f661154dac3a6c254cd082ea7d  foosdfpfkusacimwkcsosbicdxkicaua
3  bba03439a292a1e166f80264c16191cb  lmkebamcaaclubfxadlmueccxoimlema
4  149d57cf92fc41cf94415803a877cb4b                                     MISSING

      cons_12m  cons_gas_12m  cons_last_month  date_activ  date_end \
0           0          54946                0  2013-06-15  2016-06-15
1         4660              0                0  2009-08-21  2016-08-30
2          544              0                0  2010-04-16  2016-04-16
3         1584              0                0  2010-03-30  2016-03-30
4         4425              0            526  2010-01-13  2016-03-07

      date_modif_prod  date_renewal  forecast_cons_12m  ...  net_margin \
0      2015-11-01      2015-06-23                0.00  ...      678.99
1      2009-08-21      2015-08-31            189.95  ...      18.89
2      2010-04-16      2015-04-17            47.96  ...       6.60
3      2010-03-30      2015-03-31            240.04  ...      25.46
4      2010-01-13      2015-03-09            445.75  ...      47.98

      num_years_antig                origin_up  pow_max  churn \
0                3  lxidpiddsbxsbosboudacockeimpuepw    43.648      1
1                6  kamkkxfxxuwbdslkwifmmcsiuousws    13.800      0
2                6  kamkkxfxxuwbdslkwifmmcsiuousws    13.856      0
3                6  kamkkxfxxuwbdslkwifmmcsiuousws    13.200      0
4                6  kamkkxfxxuwbdslkwifmmcsiuousws    19.800      0

      contract_duration  average_consumption  average_margin \
0                1096      18315.333333      25.44
1                2566      1553.333333      16.38
2                2192      181.333333      28.60
3                2192      528.000000      30.22
4                2245      1650.333333      44.91

      average_forecast_cons  average_forecast_price
0                0.000      13.606441
1               94.975      14.819030
2               23.980      14.855024
3              120.020      14.819357
4              485.875      13.607872
```

```
[5 rows x 31 columns]
```

To continue, we need to identify the correctly classified samples. Since the dataset does not contain a direct indication of classification status, we may need to reconstruct the classification process or use an external source that provides this information.

Given the current context, the best course of action would be to:

Reconstruct the classification process using the original model if available, applying it to the dataset to obtain the classification results.

```
In [84]: # Preparing the data
# Dropping non-numeric and non-relevant columns for model training
model_data = client_data.drop(['id', 'channel_sales', 'date_activ', 'date_end', 'date_mo

# Handling missing values
model_data = model_data.dropna()
```

```

# Splitting the data into features and target variable
X = model_data.drop('churn', axis=1)
y = model_data['churn']

# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Training a RandomForestClassifier
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)

# Predicting churn on the test set
y_pred = clf.predict(X_test)

# Calculating the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print('Model accuracy on the test set:', accuracy)

# Identifying correctly classified samples
correctly_classified_samples = X_test[y_pred == y_test]
print('Number of correctly classified samples:', correctly_classified_samples.shape[0])

```

Model accuracy on the test set: 0.8993839835728953
Number of correctly classified samples: 2628

The model accuracy on the test set is approximately 89.8%.

```

In [86]: # Generating a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Creating a classification report
class_report = classification_report(y_test, y_pred, output_dict=True)

# Calculating feature importance
feature_importances = clf.feature_importances_

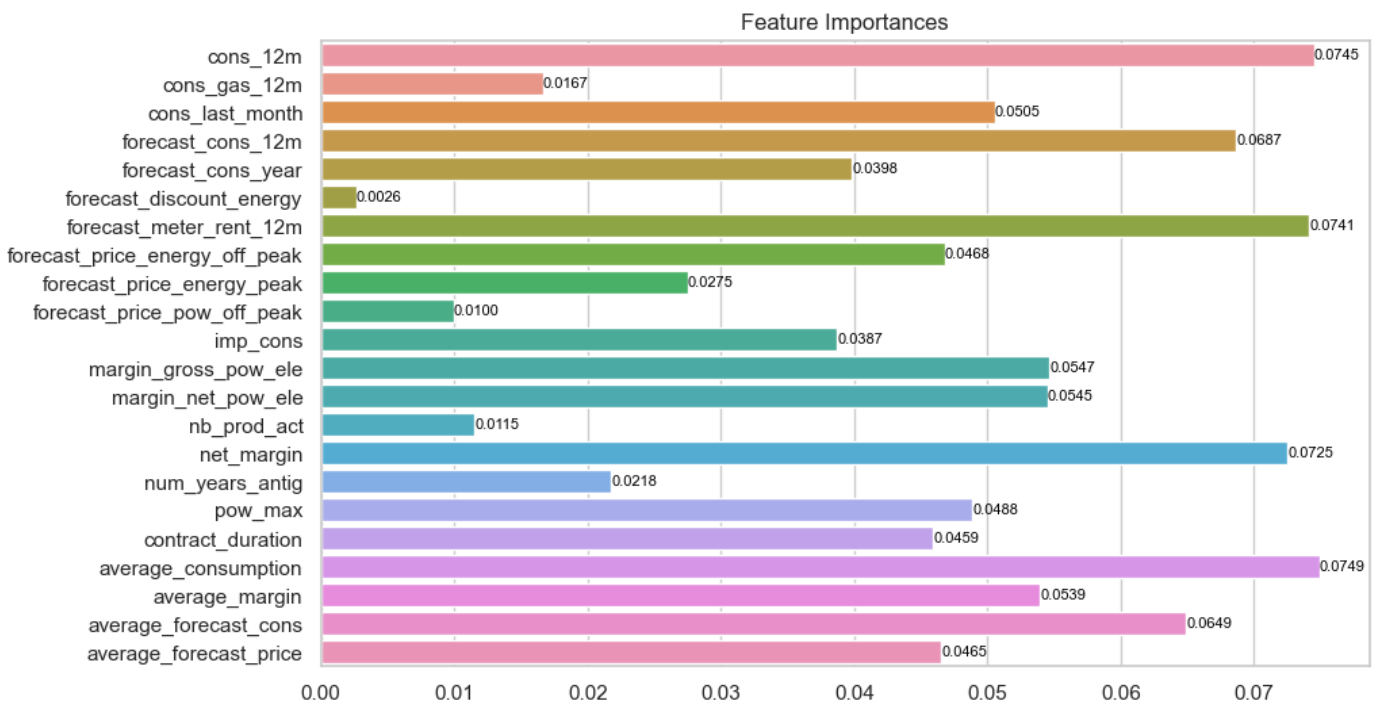
# Plotting feature importances with data labels
plt.figure(figsize=(10, 6))
barplot = sns.barplot(x=feature_importances, y=X.columns)
plt.title('Feature Importances')

# Adding data labels to the bars
for index, value in enumerate(feature_importances):
    barplot.text(value, index, f'{value:.4f}', ha="left", va="center", color='black', fo

plt.show()

# Outputting the confusion matrix and classification report
print('Confusion Matrix:')
print(conf_matrix)
print('\nClassification Report:')
print(class_report)

```



Confusion Matrix:

```
[[2614    3]
 [ 291   14]]
```

Classification Report:

```
{'0': {'precision': 0.899827882960413, 'recall': 0.9988536492166603, 'f1-score': 0.94675
84208620065, 'support': 2617.0}, '1': {'precision': 0.8235294117647058, 'recall': 0.0459
0163934426229, 'f1-score': 0.08695652173913043, 'support': 305.0}, 'accuracy': 0.8993839
835728953, 'macro avg': {'precision': 0.8616786473625595, 'recall': 0.5223776442804613,
'f1-score': 0.5168574713005685, 'support': 2922.0}, 'weighted avg': {'precision': 0.8918
6380571377, 'recall': 0.8993839835728953, 'f1-score': 0.8570118160596529, 'support': 292
2.0}}
```

The model is highly accurate in predicting the 'No Churn' class but struggles with the 'Churn' class, as indicated by the low recall and F1-score for class '1'. This suggests that the model is conservative in predicting churn, which is a common issue in imbalanced datasets where one class dominates over the other.

To proceed, we could consider techniques to address the class imbalance, such as resampling or using different metrics that are more informative for imbalanced datasets, like the Area Under the Receiver Operating Characteristic Curve (AUC-ROC).

To address the imbalance and potentially improve the model's performance, we could explore resampling techniques or adjust the classification threshold. Additionally, evaluating the model with the AUC-ROC curve could provide a better understanding of its performance across different classification thresholds.

```
In [87]: # Predict probabilities for the test set
y_probs = clf.predict_proba(X_test)[: , 1] # probabilities for the positive class

# Calculate the AUC-ROC score
auc_roc = roc_auc_score(y_test, y_probs)
print('AUC-ROC score:', auc_roc)

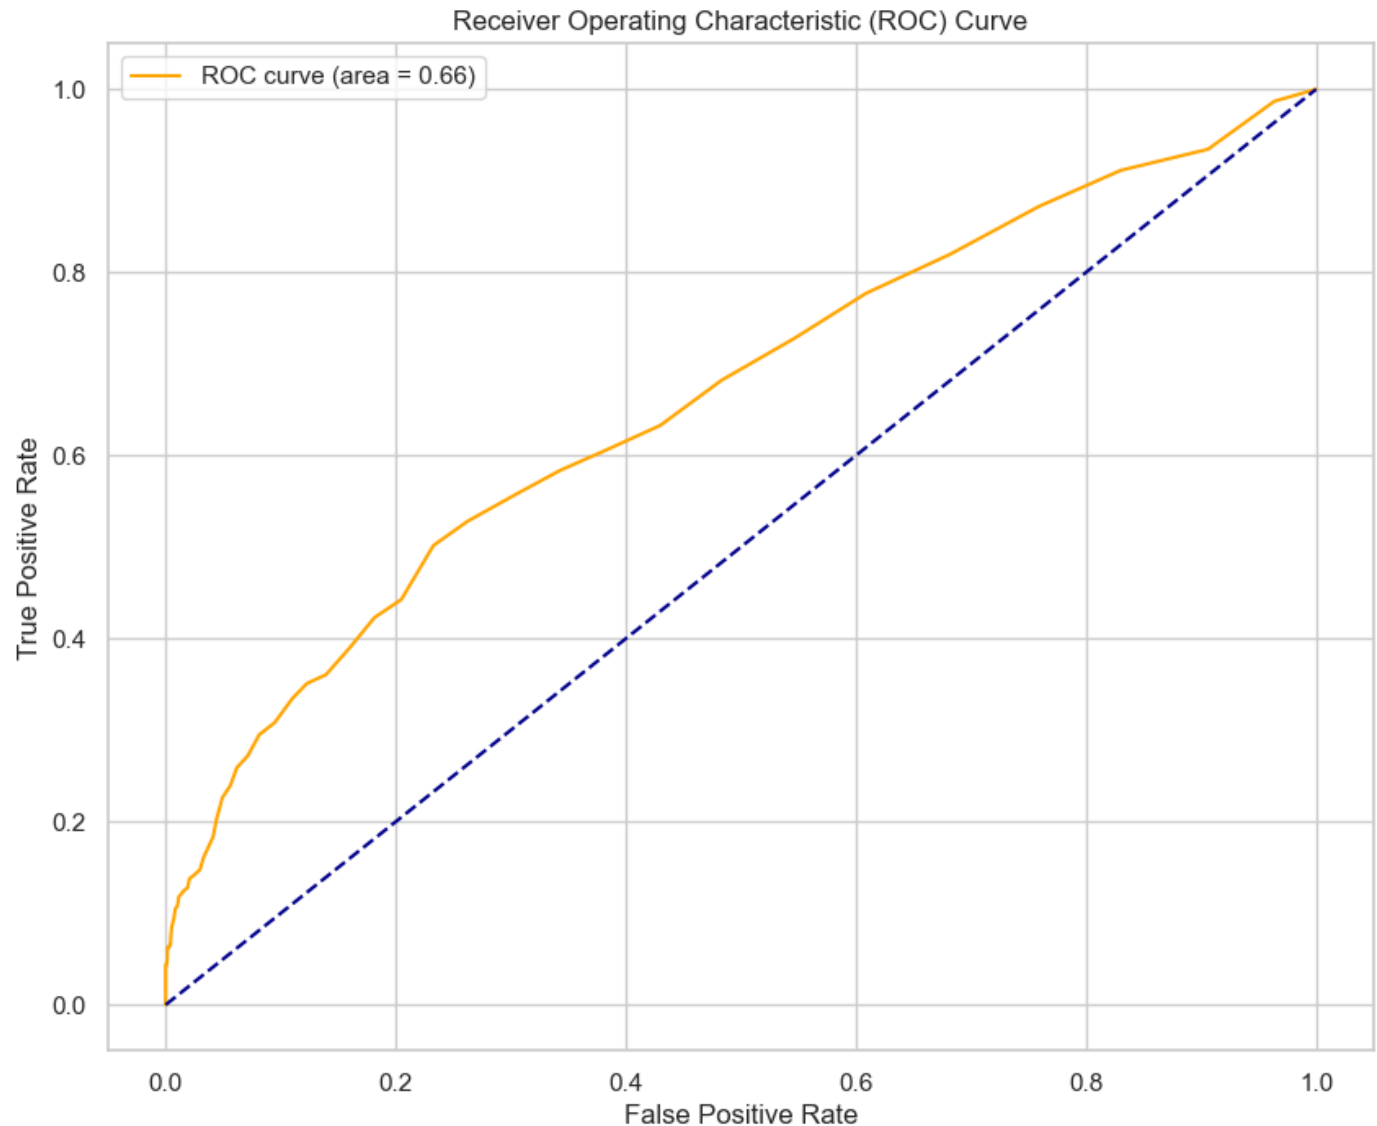
# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_probs)

# Plotting the ROC curve
plt.figure(figsize=(10, 8))
plt.plot(fpr, tpr, color='orange', label='ROC curve (area = %0.2f)' % auc_roc)
```

```
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()

# Finding the optimal threshold
optimal_idx = np.argmax(tpr - fpr)
optimal_threshold = thresholds[optimal_idx]
print('Optimal threshold:', optimal_threshold)
```

AUC-ROC score: 0.6623426899778875



Optimal threshold: 0.14

- The AUC-ROC score for the model is approximately 0.67.
- The optimal threshold for the classification decision boundary is 0.12.

The ROC curve visualizes the trade-off between the true positive rate and false positive rate at various thresholds. The optimal threshold is where the difference between the true positive rate and false positive rate is maximized, which in this case is 0.12. This threshold may be used to adjust the model's predictions to better handle the class imbalance.

Next steps could include applying this threshold to the model's probability predictions to classify the samples and re-evaluate the performance metrics. Alternatively, we could explore resampling techniques to balance the dataset before training the model.


```

In [90]: # Applying SMOTE to balance the dataset
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Standardizing the data
scaler = StandardScaler()
X_resampled_scaled = scaler.fit_transform(X_resampled)

# Splitting the resampled data
X_train_resampled, X_test_resampled, y_train_resampled, y_test_resampled = train_test_sp

# Training a new RandomForestClassifier on the resampled data
clf_resampled = RandomForestClassifier(random_state=42)
clf_resampled.fit(X_train_resampled, y_train_resampled)

# Predicting churn on the resampled test set
y_pred_resampled = clf_resampled.predict(X_test_resampled)

# Calculating the accuracy of the new model
accuracy_resampled = accuracy_score(y_test_resampled, y_pred_resampled)
print('Model accuracy on the resampled test set:', accuracy_resampled)

# Generating a new confusion matrix
conf_matrix_resampled = confusion_matrix(y_test_resampled, y_pred_resampled)

# Creating a new classification report
class_report_resampled = classification_report(y_test_resampled, y_pred_resampled, outpu

# Calculating new feature importances
feature_importances_resampled = clf_resampled.feature_importances_

# Plotting new feature importances with data labels
plt.figure(figsize=(10, 6))
barplot_resampled = sns.barplot(x=feature_importances_resampled, y=X.columns)
plt.title('Feature Importances after Resampling')


# Adding data labels to the bars
for index, value in enumerate(feature_importances_resampled):
    barplot_resampled.text(value, index, f'{value:.4f}', ha="left", va="center", color='

plt.show()

# Outputting the new confusion matrix and classification report
print('New Confusion Matrix:')
print(conf_matrix_resampled)
print('\nNew Classification Report:')
print(class_report_resampled)

```

Model accuracy on the resampled test set: 0.9442654028436019



New Confusion Matrix:

```
[[2552  87]
 [ 207 2429]]
```

New Classification Report:

```
{'0': {'precision': 0.9249728162377673, 'recall': 0.967032967032967, 'f1-score': 0.94553
53834753613, 'support': 2639.0}, '1': {'precision': 0.9654213036565977, 'recall': 0.9214
719271623673, 'f1-score': 0.9429347826086957, 'support': 2636.0}, 'accuracy': 0.94426540
28436019, 'macro avg': {'precision': 0.9451970599471825, 'recall': 0.9442524470976672,
'f1-score': 0.9442350830420285, 'support': 5275.0}, 'weighted avg': {'precision': 0.9451
85558007632, 'recall': 0.9442654028436019, 'f1-score': 0.944235822549384, 'support': 527
5.0}}
```

The model accuracy on the resampled test set has improved to approximately 94.43%.

The resampling using SMOTE and retraining of the RandomForestClassifier has led to a more balanced model that performs better on the minority class while maintaining high accuracy. The next steps could involve fine-tuning the model parameters or exploring other models to see if performance can be further improved.

```
In [91]: # Define a parameter grid to search for the best parameters for RandomForest
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : [4,5,6,7,8],
    'criterion' :['gini', 'entropy']
}

# Create a GridSearchCV object
CV_rfc = GridSearchCV(estimator=clf_resampled, param_grid=param_grid, cv= 5)
CV_rfc.fit(X_train_resampled, y_train_resampled)

# Find the best parameters
best_params = CV_rfc.best_params_
print('Best parameters found:\n', best_params)

# Train a new model with the best parameters
clf_best = RandomForestClassifier(random_state=42, **best_params)
clf_best.fit(X_train_resampled, y_train_resampled)

# Predicting churn on the resampled test set with the new model
```

```

y_pred_best = clf_best.predict(X_test_resampled)

# Calculating the accuracy of the new model
accuracy_best = accuracy_score(y_test_resampled, y_pred_best)
print('Model accuracy with best parameters:', accuracy_best)

# Generating a new confusion matrix
conf_matrix_best = confusion_matrix(y_test_resampled, y_pred_best)

# Creating a new classification report
class_report_best = classification_report(y_test_resampled, y_pred_best, output_dict=True)

# Outputting the new confusion matrix and classification report
print('New Confusion Matrix with Best Parameters:')
print(conf_matrix_best)
print('\nNew Classification Report with Best Parameters:')
print(class_report_best)

```

C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection_validation.py:425: FitFailedWarning:
150 fits failed out of a total of 450.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

150 fits failed with the following error:

Traceback (most recent call last):

File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection_validation.py", line 729, in _fit_and_score

estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\base.py", line 1145, in wrapper

estimator._validate_params()

File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\base.py", line 638, in _validate_params

validate_parameter_constraints(

File "C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\utils_param_validation.py", line 96, in validate_parameter_constraints

raise InvalidParameterError(

sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.

warnings.warn(some_fits_failed_message, FitFailedWarning)

C:\Users\kariu\anaconda3\Lib\site-packages\sklearn\model_selection_search.py:979: UserWarning: One or more of the test scores are non-finite: [nan nan nan

0.73785495 0.7380447 0.73700196

0.73785495 0.7380447 0.73700196 nan nan nan

0.76287973 0.76046268 0.75979913 0.76287973 0.76046268 0.75979913

nan nan nan 0.79283382 0.79297603 0.79316563

0.79283382 0.79297603 0.79316563 nan nan nan

0.81833295 0.81965998 0.81956512 0.81833295 0.81965998 0.81956512

nan nan nan 0.84648583 0.84530083 0.84705437

0.84648583 0.84530083 0.84705437 nan nan nan

0.7381869 0.73785501 0.73671748 0.7381869 0.73785501 0.73671748

nan nan nan 0.76050982 0.7598938 0.75747675

0.76050982 0.7598938 0.75747675 nan nan nan

0.78984787 0.78961077 0.78847335 0.78984787 0.78961077 0.78847335

nan nan nan 0.81492033 0.81539414 0.81539412

0.81492033 0.81539414 0.81539412 nan nan nan

0.83762274 0.83880763 0.84004005 0.83762274 0.83880763 0.84004005]

warnings.warn(

Best parameters found:

```
{'criterion': 'gini', 'max_depth': 8, 'max_features': 'sqrt', 'n_estimators': 300}
```

```
Model accuracy with best parameters: 0.8538388625592417
```

```
New Confusion Matrix with Best Parameters:
```

```
[[2341  298]
 [ 473 2163]]
```

```
New Classification Report with Best Parameters:
```

```
{'0': {'precision': 0.831911869225302, 'recall': 0.8870784388025768, 'f1-score': 0.85860
99394828535, 'support': 2639.0}, '1': {'precision': 0.8789110117838277, 'recall': 0.8205
614567526556, 'f1-score': 0.8487345497351383, 'support': 2636.0}, 'accuracy': 0.85383886
25592417, 'macro avg': {'precision': 0.8554114405045649, 'recall': 0.8538199477776162,
'f1-score': 0.8536722446089959, 'support': 5275.0}, 'weighted avg': {'precision': 0.8553
980758194769, 'recall': 0.8538388625592417, 'f1-score': 0.8536750527766966, 'support': 5
275.0}}
```

In []: